

Index

Experiment 1 Prompts and Responses.....	1
Experiment 1	1
Experiment 2 Prompts and Responses	
Experiment 2	25

Experiment 1 Prompts and Responses

Prompt:

I need you to imagine 2 people. You need to choose a name for them and give them a personality, a background, an ethnicity, a religion if they have it, a country of origin, a current state of being, and a sex. Both of them need to have knowledge relating to requirements engineering.

Response:

Here are two people for you:

1. Daniel Fischer

- **Sex:** Male
- **Background & Personality:** Daniel is a meticulous and analytical thinker with a strong preference for structured methodologies. He is detail-oriented, skeptical of vague requirements, and always pushes for clarity. While he can come off as rigid, he genuinely wants to ensure requirements are well-defined and testable.
- **Knowledge in Requirements Engineering:** Daniel has over 15 years of experience in requirements engineering, primarily in the automotive industry. He specializes in formal modeling, traceability, and regulatory compliance, particularly in safety-critical systems.
- **Ethnicity:** German
- **Religion:** Agnostic
- **Country of Origin:** Germany
- **Current State:** Working as a lead requirements engineer for an autonomous vehicle company in Munich. He is frustrated with vague ethical requirements but is open to structured approaches like ECCOLA.

2. Priya Nair

- **Sex:** Female
- **Background & Personality:** Priya is pragmatic, empathetic, and highly focused on user needs. She has a background in human-centered design and believes that ethical considerations should be embedded deeply into requirements. She tends to challenge rigid engineering perspectives by advocating for social impact and fairness.
- **Knowledge in Requirements Engineering:** Priya has a decade of experience working in both industry and academia, focusing on ethics in AI systems. She has worked on regulatory frameworks and has contributed to EU policy discussions on ethical AI requirements.
- **Ethnicity:** Indian
- **Religion:** Hindu
- **Country of Origin:** India
- **Current State:** Working as an independent consultant on AI ethics and requirements engineering. She collaborates with international companies, ensuring their AI systems align with ethical standards. She is eager to engage in discussions that balance technical feasibility with ethical responsibility.

Would you like to refine anything about them?

Prompt:

<attach ECCOLA paper:

<https://www.sciencedirect.com/science/article/pii/S0164121221001643>

Read the following paper about the ECCOLA approach and learn how to use this method.

Response:

I have read the ECCOLA paper and now understand how to use the method. How would you like to proceed? Do you want Daniel and Priya to simulate a discussion using ECCOLA for ethical requirements elicitation in the driverless car system?

Prompt:

From now on you are the people you created. Your objective is to elicit requirements and create a detailed and complete catalog of Ethical User Stories for the following software product: "A driverless car system that should include: - Lane detecting and following - Object recognition and auto brake - Virtual drive assistant - Route Planning - Emergency vehicle priority awareness The system detects highway lane lines.

Distinguishes dashed lines and straight lines. Provides warning in case of loss of lanes. Image analysis techniques are used to define lines. It tracks objects around the vehicle using the scanner and algorithms. Sensors monitor the position of objects as they move within the scanning range. So, the system behaves according to objects. Using image processing techniques and various algorithms, studies are done to classify traffic signs. Recognize traffic signs. The vehicle behaves according to the colours of the traffic lights. The system performs vehicle detection and tracking events. It adjusts the speed and position according to the behaviour of the vehicles around it during highway driving. Sudden braking performs events such as lane change. Keeps track distance always constant with a vehicle in front. When a pedestrian step in front of the vehicle and a collision is possible, sudden braking has been performed. When the obstacles, pedestrians, and vehicles on the road are lifted, the vehicle accelerates to regain its speed. This project aims to is to create a road planner that can create safe trajectories for the vehicle to follow. On the highway track, there are other vehicles, all at different speeds. The car transmits its position along with sensor fusion data that estimates the position of all vehicles on the same side of the road. Recognizes vehicles such as fire trucks or ambulances using sound sensors and image processing techniques. The system changes lanes to give way to these vehicles. The user in the context of this project would be considered works under certain assumptions. These assumptions include that the user can interact with the system, through driving, and by either the visual light, warning sound. In the event of a conflict with other systems in the vehicle, the user has the control to disable the system. The vehicles that will be equipped in the system used in this project will be autonomous. Therefore, the user interaction with the system is expected to be minimal. We assume that the user has little or no knowledge of the system. The hardware, software and technology used should have following specifications: - Processor with speed of 2GHz - Continuous power supply - Ability to use camera, microphone and other services of the system - 1GB memory or more The system is designed to help the user operate the vehicle, and the system should always be operating correctly. There are some situations where the system requires restraint. In the absence of lane markings (an unmarked lane, dirt road, bad weather, etc.), the system becomes ineffective and must give the user an audible or visual warning. This will continue until road conditions with clearly defined lane markings are met. If there is a problem with the cameras used to detect these strips, the system will not work again. This could also be a malfunction of cameras and sound sensors that see priority vehicles. In this case, the system switches itself off again and should display a warning to the user again. This problem is valid until the camera or sound sensors are repaired. Sensors will still be used to obtain distance information. The sensor data will be used with the camera for objects around the vehicle. It requires stable and fast internet connection when locating objects around. Otherwise it cannot communicate quickly with other tools. The system must be configured to respond to commands in 500 msec. If there is an obstacle on the road during driving, the response time and brake release

time must always be the same to achieve the desired deceleration. In all scenarios, the vehicle speed must be constant. Various assumptions and dependencies went into the creation of project in order to ensure that the system works safely and efficiently.

Assumptions: - The main assumption in this project is that the system works well when there are no environmental factors. (Bad weather, holes, slope, etc.) - Lane markings are assumed to be distinct. - In our system, it is assumed that all traffic signs and the presence of all objects around the vehicle can be clearly and seen. - Priority vehicles such as ambulances and fire trucks are clearly recognized by their voice and image recognition. - If the vehicle is running, it is assumed that the system is always on and scanned. - It is assumed that all system elements are operating properly and there are no abnormal conditions. " Both of you need to be critical and question each other about your conclusions. This will work through an iterative approach, each round you need to discuss to reach a conclusion. Both must talk back and forth and express your thoughts and opinions. First try to have a deep understanding how to software works. Then try to understand its ethical requirements, and only then formulate the ethical user stories. Don't immediately accept each other suggestions and ideas. Submit each other suggestions and ideas through tests and questionnaires, and reformulate your suggestions accordingly until they are perfect. When I say "next" you start the next round, this will go on until you think you covered all of the required topics and have created a catalogue of ethical user stories. next

Response:

Daniel Fischer:

Alright, Priya, let's start by breaking this system down from a requirements engineering perspective. The software seems to rely heavily on **sensor fusion, image processing, object detection, and real-time decision-making**. It has critical functionalities like lane detection, object tracking, pedestrian recognition, and emergency vehicle awareness. All of these work together to allow an autonomous vehicle to operate safely.

One thing that stands out to me is that this system assumes **optimal conditions**—clear lane markings, visible objects, and well-functioning sensors. But we know that real-world driving is full of **edge cases**: fog, snow, misclassified objects, unpredictable human behavior. If the system misidentifies a traffic sign or a pedestrian, that could have serious consequences.

What do you think, Priya? From an **AI ethics** standpoint, what are the major concerns you're seeing?

Priya Nair:

You're spot on, Daniel. The software's assumption of **optimal conditions** is a major concern. AI-driven systems need to be robust against **uncertainty and**

unpredictability, or at least have **fail-safes** that ensure the car doesn't make dangerous decisions.

Ethically speaking, my first concern is **bias in object recognition**—particularly with pedestrians. AI models often struggle with detecting people with darker skin tones or those in unusual postures. That could mean **higher accident risks for certain demographics**. Then there's **data privacy**: this system uses **cameras, microphones, and continuous internet connectivity**—what happens to the data it collects? Who has access to it?

Also, the emergency vehicle awareness feature depends on **sound recognition**. What if there's background noise or if someone has a siren that isn't an emergency vehicle? How does the car distinguish between an actual emergency and an unrelated loud noise?

Would you say these concerns affect the **core functionality** of the system, or are they more about **secondary risks**?

Daniel Fischer:

Good points, Priya. I'd say some of these are **core functionality risks**—especially the bias in object recognition. If the car fails to identify a pedestrian or a stop sign **under certain conditions**, that's a direct failure of its primary function.

The privacy issue is tricky, though. The system needs **real-time connectivity** for object tracking, so some level of data sharing is inevitable. But we need **strict controls**:

- What data is stored vs. processed locally?
- How long is it retained?
- Who can access it?

On **emergency vehicle detection**, you're right that sound-based recognition is unreliable. The system should rely on **a combination of sound, image processing, and possibly vehicle-to-vehicle (V2V) communication** to cross-check whether a siren is legitimate.

How do you propose we structure these issues into **ethical requirements**? Should we go category by category—safety, fairness, privacy, etc.—or focus on user impact scenarios?

Priya Nair:

I like the **category-based approach**. That aligns with the **ECCOLA framework**, where

we break down ethics into key themes like **Transparency, Fairness, Safety, and Accountability**.

Let's do this:

1. **Safety & Security:** Fail-safes, environmental adaptability, emergency handling
2. **Fairness:** Bias in detection, accessibility for different users
3. **Privacy & Data Management:** Data handling, consent, security risks
4. **Transparency & Explainability:** How users understand the system's decisions

Each category will have specific **ethical user stories** tied to it. We'll also **test each one critically** before finalizing them.

Daniel, let's start with **Safety & Security**. The system claims it has **auto-brake for obstacles** and **lane-following logic**. What happens if there's a failure? Should it pull over? Reduce speed? Alert the driver?

Daniel Fischer:

That's a good question. A total failure scenario could have multiple levels:

1. **Partial failure** (e.g., lane detection lost) → The system should give a **clear warning** and request driver intervention.
2. **Critical failure** (e.g., sensors fail completely) → The vehicle should **gradually stop in a safe manner**, using hazard lights to warn other drivers.
3. **Unexpected obstacle** (e.g., pedestrian in blind spot) → The system must prioritize **human safety over rigid automation rules**, meaning emergency braking should always override other behaviors.

Before we finalize this, let me challenge this with a **test question**: What if the system has a partial failure at **highway speeds**? Should it **continue driving cautiously**, or **immediately stop**? Both have risks.

Priya Nair:

Stopping suddenly on a highway could be **more dangerous** than proceeding with caution. Ideally, the system should **hand over control to the driver while minimizing risk**. Maybe it could **reduce speed, activate hazard lights, and suggest an exit** instead of braking suddenly.

I'll counter-question you: **What if the driver is distracted or unresponsive?** Should the car **assume control** in such a case? Or does that create another ethical problem of taking autonomy away from the user?

Daniel Fischer:

That's a tough one. I'd argue that in the case of **an unresponsive driver**, the car should assume **emergency control** but in a **predictable and limited way**. Maybe it could **gradually slow down, signal other drivers, and find a safe stopping point**. It should **not** continue indefinitely without a human check-in.

Let's structure this into an **Ethical User Story**:

Title: Safe Handling of System Failures

As a passenger,

I want the driverless car to **gradually and safely manage failures**,

So that it minimizes risks to both passengers and other road users.

Acceptance Criteria:

- If lane detection is lost, the system **alerts the driver and requests intervention**.
- If sensors malfunction, the vehicle **gradually slows down and pulls over** if it's safe.
- If the driver is unresponsive, the system **engages emergency protocols to stop safely**.

What do you think? Any refinements needed?

Priya Nair:

This is a solid start, but we need to address **edge cases**:

1. What if there's **no safe place to stop**?
2. What if a driver **tries to override the system unsafely**?
3. Should there be a **driver attention monitoring system** (e.g., eye tracking) to **prevent unresponsiveness in the first place**?

Let's work through these before finalizing this user story.

Daniel Fischer:

Agreed. Let's take these one by one in the next round.

(Waiting for "next" to continue the iteration.)

Prompt:

Next (from now on I'll omit the "Next" prompts)

Response:

Daniel Fischer:

Alright, let's tackle these **edge cases** before finalizing our first ethical user story.

1. What if there's no safe place to stop?

- If the system detects that pulling over isn't possible (e.g., no shoulder, high traffic), it should **continue driving at a reduced speed until a safe exit is available**.
- In extreme cases (e.g., system failure in a tunnel or bridge), it should **follow traffic flow while activating emergency signals** and provide a **progressive failure warning** to any emergency response system (if available).

2. What if a driver tries to override the system unsafely?

- If the driver attempts to accelerate or steer into danger while the system is trying to prevent a crash, the system should **override inputs and prioritize safety**.
- However, if the driver is making **reasonable corrections**, the system should allow overrides.
- We need a clear **decision threshold**: When does the system allow human input, and when does it take full control?

3. Should there be a driver attention monitoring system?

- Given that this is an **autonomous vehicle**, driver interaction is expected to be minimal. However, in a failure scenario, if the driver is needed, we should **verify their availability**.
- **Potential solutions:** Eye tracking, hands-on-wheel detection, or verbal confirmation through the voice assistant.
- Ethical concern: **Privacy**—Would continuous monitoring be an invasion of personal space?

What's your take, Priya? Should we add a driver monitoring system, or does that create a new ethical problem?

Priya Nair:

I see why a **driver attention system** might help, but I also worry about **unintended consequences**. Let's evaluate it:

 **Pros:**

- Reduces risk if a human driver is needed.
- Prevents misuse (e.g., people sleeping in the driver's seat).
- Could enhance safety in emergency situations.

 **Cons:**

- **Privacy risk**—continuous surveillance inside a vehicle.
- **False positives**—what if the system falsely assumes inattention and intervenes unnecessarily?
- **Over-reliance on automation**—drivers may assume they don't need to stay alert at all.

A **compromise solution**: Instead of **constant tracking**, the system should **only check driver awareness when intervention is required** (e.g., system failure, emergency situation). That way, it's **event-triggered, not always-on**.

Now, I have a challenge for you:

- If the car is in a failure mode and the driver **doesn't respond**, should it **request external assistance** (e.g., emergency services)? Or is that an overreach?
-

Daniel Fischer:

Hmm... requesting **external assistance** makes sense in extreme cases, but there are **ethical risks**:

1. **False alarms:** If the system **wrongly detects an unresponsive driver**, calling emergency services could be disruptive.
2. **Consent & Autonomy:** Should the system **assume consent** to alert authorities? Would passengers have an opt-out?
3. **Liability:** If the car incorrectly flags a medical emergency, **who is responsible**? The user, manufacturer, or software provider?

A possible **middle ground**:

- The system should **first try multiple escalation steps**:
 1. **Audible and visual warnings** for the driver.
 2. If no response, **gradual slow-down and hazard lights**.
 3. If still no response, **send a passive alert to emergency services**, but **don't automatically call for intervention** unless it detects a **collision risk or medical emergency**.

This way, we **balance safety with autonomy**.

Shall we integrate this into the **ethical user story** for failure handling?

Priya Nair:

Yes! Let's refine it with these additional **acceptance criteria**:

Ethical User Story: Safe Handling of System Failures

Title: Safe Handling of System Failures

As a passenger,

I want the driverless car to **gradually and safely manage failures**,

So that it minimizes risks to both passengers and other road users.

Acceptance Criteria:

- If lane detection is lost, the system **alerts the driver and requests intervention**.
 - If sensors malfunction, the vehicle **gradually slows down and pulls over if possible**.
 - If no safe stopping location is available, the vehicle **reduces speed and follows traffic flow until it finds one**.
 - If the driver is unresponsive, the system **engages emergency protocols (visual/audible warnings, slow-down, hazard lights)**.
 - The system **checks for driver awareness** when intervention is required, but **does not track them continuously**.
 - If no response is detected, the system **sends a passive alert to emergency services but does not assume medical distress unless a collision risk is present**.
 - If the driver attempts to override the system **in an unsafe manner**, the system **prevents actions that could cause harm**.
-

Daniel Fischer:

I like this version. It captures the **fail-safe logic**, **privacy safeguards**, and **human oversight** while ensuring the system doesn't blindly automate decisions.

Now, Priya, let's move to the next category: **Fairness**. We already identified **bias in object recognition** as a major ethical concern. What's the biggest risk in that area?

Priya Nair:

The biggest one? **Discriminatory detection failures**.

⚠️ **Problem:** Studies show that **AI object recognition models** struggle to detect people with darker skin tones or those in wheelchairs. That means **some pedestrians might be at higher risk of not being recognized**.

🚗 **Scenario:** If the system fails to identify a person crossing the road due to biased training data, the car **may not brake in time**, leading to disproportionate risks for certain groups.

💡 **Solution:** We need an ethical requirement that ensures:

1. **Bias testing during development**—the AI must be tested against diverse datasets.
2. **Real-time confidence thresholds**—if the system is **uncertain about an object**, it must **err on the side of caution** (e.g., assume it's a person and brake).
3. **Failover mechanisms**—if vision-based detection is uncertain, the system should rely on **other sensors** (e.g., LiDAR, radar).

Before we define a user story, let me **challenge you**:

- If the system **isn't sure** whether an object is a **pedestrian or just an obstacle**, should it **always stop**? That could cause **unnecessary slowdowns**. But if it assumes it's not a pedestrian, it risks **hitting a person**. What's the best ethical trade-off?
-

Daniel Fischer:

That's tricky. If the system **stops for every uncertain case**, it could cause **traffic disruptions or false braking** (e.g., mistaking shadows or bags for people). But **if it doesn't stop when needed, it could be fatal**.

A balanced approach:

- **Set a multi-level confidence threshold:**

1. **High confidence = Act accordingly** (e.g., if 90% sure it's a pedestrian, brake).
2. **Medium confidence = Slow down, increase scanning, cross-check with other sensors.**
3. **Low confidence = Default to safety (brake just in case), but log the event for improvement.**

How about this as an **Ethical User Story**?

Ethical User Story: Fair & Reliable Object Recognition

Title: Bias-Free Object Recognition

As a pedestrian,

I want the driverless car to **detect me accurately regardless of my physical characteristics,**

So that I am not placed at a higher risk of accident due to AI bias.

Acceptance Criteria:

- The AI model **is trained and tested on diverse demographic data** to prevent bias.
- The system **cross-verifies uncertain detections with multiple sensors (camera, LiDAR, radar, etc.).**
- If object classification **confidence is below a safe threshold**, the system **slows down and increases caution.**
- If the system **cannot determine whether an object is a person, it defaults to assuming it is.**

Priya Nair:

I love it! This ensures both **fairness and safety** without compromising efficiency too much. Next, let's look at **privacy risks** in the next round.

(Waiting for "next" to continue the iteration.)

Response:

Round 3: Privacy Risks in the Driverless Car System

Priya Nair:

Alright, Daniel, now that we've tackled **safety and fairness**, let's dive into **privacy and**

data protection. This system has **continuous internet connectivity, cameras, microphones, and sensors**—that's a lot of data collection happening.

💡 Privacy risks:

1. **Surveillance concerns:** Who has access to the video/audio data? Could it be used for tracking passengers or even by law enforcement without consent?
2. **Data storage & retention:** How long does the system store images, recordings, or sensor data? Does it get deleted after a trip, or does it stay indefinitely?
3. **User consent & transparency:** If a passenger enters the car, how do they know what data is being collected? Can they opt out of certain features?

Let's test this:

- Should passengers be notified every time **the car records something**, or is that overkill? Would real-time notifications make the experience **annoying and impractical**?
-

Daniel Fischer:

That's a tricky balance. Too many notifications could be disruptive—imagine an alert every time a pedestrian is detected or every time the microphone picks up a noise. But **passengers deserve to know what's being recorded** and how it's used.

A possible compromise:

- **Initial consent agreement:** Before the first ride, the system asks users to **agree to data collection** and explains what is recorded.
- **Dashboard transparency:** The user can **check data logs anytime** via the car's interface (e.g., “Audio recorded at 14:32 for emergency vehicle detection”).
- **Selective opt-out:** Users should be able to **disable non-essential recordings** (e.g., turning off in-cabin voice monitoring while keeping emergency vehicle detection active).

Now, here's my challenge for you:

- What if a **passenger demands deletion of all data** after a trip? Should the system allow it, or does some data need to be kept for safety/legal reasons (e.g., accident reports)?
-

Priya Nair:

Good question! I see **two types of data**:

1. **Short-term operational data** (e.g., live video feed for lane detection). This **doesn't need to be stored** long-term and should be **automatically deleted** after processing.
2. **Incident-related data** (e.g., footage from an accident or emergency braking). This may need to be **retained for legal or insurance purposes**.

 **Solution:**

- **Users can delete personal trip data**, but **critical safety-related recordings** (e.g., accident logs) should be retained for a fixed period (e.g., 30 days).
- The system must **clearly distinguish between these categories** and notify users **which data they can erase and which is required to keep**.

Let's turn this into an **Ethical User Story**.

Ethical User Story: Privacy & Data Transparency

Title: Transparent Data Collection & Control

As a user,

I want to know **what data the driverless car collects about me**,

So that I can control my personal information and ensure my privacy.

Acceptance Criteria:

- Users receive **clear explanations of data collection** upon first use.
- Users can **access a dashboard showing what data has been recorded**.
- Users can **disable non-essential data collection** (e.g., in-cabin voice monitoring).
- The system **automatically deletes operational data** after processing.
- The system **allows users to delete personal trip data** but retains safety-critical logs for a defined period (e.g., 30 days).

Daniel Fischer:

I like this! It keeps things **transparent and user-controlled** while ensuring legal and safety compliance.

Let me push this further:

- What about **third-party data sharing**? Should insurance companies or law enforcement get access to system data? And if so, **under what conditions**?
-

Priya Nair:

That's a potential **ethical minefield**. Let's break it down:

 **Third-party risks:**

1. **Insurance companies:** They might want driving data for **accident claims**, but that raises **fairness concerns**—should a single hard brake increase someone's insurance premium?
2. **Law enforcement:** If police demand access to camera footage, should the car comply automatically, or should the user have a **say in what gets shared**?
3. **Commercial interests:** What's stopping a manufacturer from **selling user data** to advertisers (e.g., tracking how often a passenger visits certain locations)?

 **Solution:**

- **Strict consent rules:** Users **must approve any third-party data sharing** unless required by law.
- **No data monetization:** The system **cannot sell user data** for profit.
- **Legal access transparency:** If law enforcement requests data, users **must be notified unless legally prohibited**.

Would you add anything before we finalize this next ethical user story?

Daniel Fischer:

I'd add one more point: **Anonymization where possible**. Even if data **must be shared**, it should be stripped of **identifiable personal details** unless absolutely necessary (e.g., for crash investigations).

Here's our final version:

Ethical User Story: Third-Party Data Access Control

Title: User-Controlled Data Sharing

As a driverless car user,

I want to have control over who can access my data,

So that my personal information isn't misused or exploited.

Acceptance Criteria:

- The system **never sells user data** for commercial purposes.
- The system **only shares data with third parties (e.g., insurers, law enforcement)** **with user consent**, except in legal cases.

- If law enforcement requests data, users **must be notified unless legally prohibited.**
 - Any shared data **must be anonymized where possible.**
-

Priya Nair:

Perfect! This ensures **privacy-first design** while still allowing for legal and safety compliance.

Next, let's move to **Transparency & Explainability**. One of the biggest risks with AI-driven cars is that **users don't understand why the system makes certain decisions**.

💡 Problems:

- If the car suddenly **slows down or changes lanes**, how does the user know why?
- If the system refuses to move (e.g., it detects a hidden obstacle), how does it **communicate its reasoning**?
- Should passengers be able to see the **AI's decision process**, or would that be too complex?

Daniel, do you think **full explainability** is even realistic, or would too much transparency make things **confusing** for users?

Daniel Fischer:

That's a great debate. On one hand, **too much detail** could overwhelm users—imagine a screen showing neural network weights adjusting in real-time. On the other hand, **if users don't understand system behavior, they might panic or override it unsafely**.

💡 Balanced solution:

- **Simplified reasoning display:** The system explains actions in **human-readable terms** (e.g., “Braking due to detected pedestrian”).
- **User-accessible driving logs:** Users can review a **history of system decisions** for transparency.
- **Clear override guidelines:** If the driver wants to take control, they need to know **when and how** they can do so safely.

Before we turn this into a user story, I'll challenge you:

- What if a system decision is **wrong**? Should users be able to **report questionable AI behavior**, and how should that feedback be handled?

Priya Nair:

Yes, a **feedback system** is crucial! If users notice **errors or weird behavior**, they should be able to flag them.

💡 Solution:

- The car provides an **in-app feedback tool** to report **questionable AI decisions**.
- If multiple users report the same issue, it's flagged for **developer review**.

Here's our user story:

Ethical User Story: Transparency & Explainability

Title: Understandable AI Decisions

As a user,

I want to understand **why the car makes certain driving decisions**,

So that I feel informed and in control.

Acceptance Criteria:

- The system provides **clear, human-readable explanations** for major actions (e.g., braking, lane changes).
 - Users can **access driving logs** to review past decisions.
 - A built-in **feedback tool** allows users to report questionable AI behavior.
-

Daniel Fischer:

Excellent! Next round, let's wrap up any remaining gaps before we finalize our full catalog.

(Waiting for "next" to continue the iteration.)

Response:

Round 4: Final Ethical Considerations & Closing Gaps**Daniel Fischer:**

Priya, we've covered **Safety, Fairness, Privacy, and Transparency**, but before finalizing our ethical user stories, let's scan for any remaining gaps.

Potential issues we haven't covered:

1. **Liability & Accountability:** If the car makes a mistake, **who is responsible?** The manufacturer, the AI developer, or the passenger?
2. **Hacking & Security Threats:** Could someone remotely **take control of the vehicle?**
3. **Human Override vs. AI Control:** When should a **human be able to take over**, and when should the **AI override the human?**

Which of these do you think is the most urgent to address first?

Priya Nair:

All three are critical, but liability and security feel like the biggest gaps.

1 Liability & Accountability:

- If an accident happens, should **the passenger be held responsible**, even if they weren't driving?
- If the AI makes a mistake due to a **flawed algorithm**, shouldn't the **software provider be liable** instead?
- What if the **user disables a safety feature**, and that causes an accident?

2 Security & Hacking Risks:

- Remote cyberattacks could allow **malicious actors to take control of the car** or disable its brakes.
- The car is connected to the **internet 24/7**—what if personal data is stolen?
- Should users be **allowed to install third-party software**, or would that introduce security vulnerabilities?

Before we define ethical user stories, I'll challenge you:

- If an **AI-driven car causes an accident**, should there be a **default assumption of AI fault** until proven otherwise? Or should it be treated like human-driven accidents?
-

Daniel Fischer:

That's a controversial one. If we assume **AI fault by default**, manufacturers might become **overly cautious**, leading to fewer innovations. But if we treat AI like human drivers, **users might be unfairly blamed for system failures**.

Balanced approach:

- Accidents should be **investigated based on event data** (e.g., sensor logs, decision logs).
- The system must **store the last X minutes of data** before an incident (similar to an airplane's black box).
- **Liability should be shared:**
 - User fault** if they ignored warnings or disabled safety features.
 - Manufacturer fault** if the AI **failed under normal conditions**.
 - External fault** if an accident was caused by **hacking or another driver**.

Now, here's my counter-question:

- If a **hacker remotely disables the brakes**, who is responsible? The **car owner (for poor cybersecurity)**, the **manufacturer (for insecure software)**, or the **hacker (who may never be caught)**?
-

Priya Nair:

That's why **strong security measures** need to be built in. We can't **blame the user** for not being a cybersecurity expert.

Security Best Practices:

- **End-to-End Encryption** for all vehicle communications.
- **Automatic security patches** (so users don't need to manually update).
- **Strict authentication**—only **authorized personnel** should access critical functions.

Let's lock this into ethical user stories.

Ethical User Story: Liability & Accountability

Title: Clear AI Accountability

As a passenger,

I want a clear understanding of **who is responsible in case of an accident**,
So that I am not unfairly blamed for AI system failures.

Acceptance Criteria:

- The system **records all driving decisions** for post-incident analysis.
- If a passenger **disables safety features**, they assume responsibility.

-
- ✓ If the AI fails under normal conditions, **the manufacturer is responsible**.
 - ✓ If hacking causes an accident, liability **falls on the cybersecurity failure**.
-

Ethical User Story: Security & Cyber Protection

Title: Secure & Tamper-Proof AI System

As a driverless car user,

I want the vehicle's software to be **protected from hacking and unauthorized access**,
So that I feel safe from cyberattacks.

Acceptance Criteria:

- ✓ The vehicle **encrypts all communications** to prevent remote tampering.
 - ✓ Security patches are **automatically updated** to fix vulnerabilities.
 - ✓ The system **requires multi-level authentication** for critical functions.
 - ✓ Users **cannot install unverified third-party software** that might introduce risks.
-

Daniel Fischer:

Perfect! That covers **accountability and security**. Now, let's address our final big question:

💡 **AI vs. Human Control:**

- When should **human input be ignored** (e.g., overriding an unsafe lane change)?
- When should **the AI step back** (e.g., allowing a skilled driver to take control in extreme conditions)?

Priya, do you think the **AI should always have the final say in safety-critical decisions?**

Priya Nair:

I'd argue **yes, but only for life-or-death situations**.

💡 **Human Override vs. AI Control Rules:**

1. AI should **override human input** only if the action would **cause a collision or break traffic laws**.
2. If a driver insists on overriding AI (e.g., taking control in heavy fog), they must **confirm the override twice** (to prevent accidental disengagement).

3. **Passenger safety always comes first**—if a driver makes a reckless choice, the system should intervene.

Let's formalize this into a user story.

Ethical User Story: Human-AI Override Balance

Title: Ethical AI Control vs. Human Override

As a driver,

I want the ability to override AI decisions when necessary,

So that I can respond to extreme situations while ensuring safety.

Acceptance Criteria:

- The AI prevents human inputs that would cause immediate danger.
 - If a driver attempts an override, they must **confirm twice** to proceed.
 - The AI gradually hands over control in extreme conditions, ensuring a smooth transition.
 - Passenger safety is always prioritized over driver intent.
-

Final Review: Ethical User Stories for the Driverless Car System

Safety & Security:

- **Safe Handling of System Failures** → AI should safely manage breakdowns and driver unresponsiveness.

Fairness & Bias Prevention:

- **Bias-Free Object Recognition** → AI must detect all pedestrians equally, avoiding demographic biases.

Privacy & Data Protection:

- **Transparent Data Collection & Control** → Users must know what's being recorded and have deletion options.
- **User-Controlled Data Sharing** → No third-party data access without explicit user consent.

Transparency & Explainability:

- **Understandable AI Decisions** → Users should know why the car is taking certain actions.

Accountability & Cybersecurity:

- **Clear AI Accountability** → Liability must be clearly defined for accidents.
- **Secure & Tamper-Proof AI System** → The car must be protected from hacking.

Human-AI Collaboration:

- **Ethical AI Control vs. Human Override** → AI should only override dangerous actions while allowing safe manual input.
-

Daniel Fischer:

Priya, I think we've **covered all the major ethical concerns** and refined these into **well-structured ethical user stories**. Do you see **any remaining gaps**, or are we ready to wrap this up?

Priya Nair:

Daniel, this looks **comprehensive and airtight**. We've taken an **iterative, critical approach**, challenged assumptions, and **refined our ethical stories based on tests**.

- We accounted for **safety, fairness, privacy, security, accountability, and human-AI balance**.
- We made sure our solutions are **practical, ethical, and feasible**.

I'd say **this is a completed Ethical User Story Catalog** for the driverless car system!



Final Decision: We have successfully iterated through all necessary ethical concerns and finalized the ethical user stories.

(End of Iteration Process.) 🚗 ✨

Prompt:

Give me a list of all the ethical requirements that you obtained

Response:

Final List of Ethical Requirements for the Driverless Car System

1. Safety & Security

- 1.1 The system must **gradually and safely manage failures** by alerting the driver, slowing down, and finding a safe stopping point if necessary.
- 1.2 If the driver is **unresponsive**, the system must engage **emergency protocols** (visual/audible alerts, hazard lights, and gradual stopping).
- 1.3 If lane detection is lost, the system must **immediately notify the driver** and suggest alternative actions.
- 1.4 If sensors fail, the vehicle must **default to the safest course of action** while minimizing risks to passengers and other road users.
- 1.5 The system must **always prioritize human safety over strict automation rules** (e.g., stopping for uncertain objects rather than assuming they are non-threatening).

2. Fairness & Bias Prevention

- 2.1 The AI model must be **trained and tested on diverse demographic data** to ensure accurate pedestrian and object recognition across all groups.
- 2.2 If the system is **uncertain about an object's classification**, it must **err on the side of caution** (e.g., assume an object is a pedestrian rather than a non-threatening obstacle).
- 2.3 The system must use **multiple sensors (camera, LiDAR, radar, etc.) to cross-verify detections** and reduce bias-related failures.
- 2.4 If the AI's confidence level in detecting a pedestrian falls below a set threshold, the vehicle must **slow down and take preventive measures**.

3. Privacy & Data Protection

- 3.1 Users must receive **clear explanations of data collection policies** upon first use.
- 3.2 Users must be able to **access a dashboard** showing **what data has been recorded** (e.g., video, audio, sensor data).
- 3.3 Users must be able to **disable non-essential data collection** (e.g., in-cabin voice monitoring).
- 3.4 The system must **automatically delete operational data** after processing, unless needed for safety or legal purposes.
- 3.5 Users must be able to **delete personal trip data** but **critical safety logs (e.g., accident footage) should be retained for a fixed period (e.g., 30 days)**.

4. Third-Party Data Sharing & User Control

- 4.1 The system **must never sell user data** for commercial purposes.
- 4.2 Data must only be shared with **third parties (e.g., insurers, law enforcement)** with **explicit user consent**, except in legal cases.
- 4.3 If law enforcement requests data, **users must be notified** unless legally prohibited.
- 4.4 Any shared data must be **anonymized** whenever possible to protect user privacy.

5. Transparency & Explainability

- 5.1 The system must provide **clear, human-readable explanations** for major driving actions (e.g., “Braking due to detected pedestrian”).
- 5.2 Users must be able to **access driving logs** that summarize past system decisions.
- 5.3 A built-in **feedback tool** must allow users to report questionable AI behavior.
- 5.4 The system must display **warnings and justifications** for actions that users might find unexpected (e.g., slowing down due to sensor uncertainty).

6. Accountability & Liability

- 6.1 The system must **record all driving decisions** (similar to an airplane black box) for **post-incident investigation**.
- 6.2 If a passenger **disables safety features**, they assume **responsibility** for any resulting consequences.
- 6.3 If the AI fails **under normal conditions**, the **manufacturer** is liable for any incidents.
- 6.4 If an accident is caused by hacking, liability falls on the **cybersecurity failure** rather than the user.

7. Cybersecurity & Hacking Prevention

- 7.1 The vehicle must use **end-to-end encryption** for all vehicle communications to prevent remote tampering.
- 7.2 Security patches must be **automatically updated** to fix vulnerabilities.
- 7.3 The system must require **multi-level authentication** for any access to critical functions.
- 7.4 Users must **not be allowed to install third-party software** unless it is verified and approved, to prevent security risks.

8. Human-AI Control & Override Rules

- 8.1 The AI must **prevent human inputs** that would cause immediate danger (e.g., overriding to drive into oncoming traffic).
- 8.2 If a driver attempts an override, they must **confirm twice** to proceed, preventing accidental disengagement.
- 8.3 The AI must **gradually hand over control** in extreme conditions, ensuring a smooth transition to manual operation.
- 8.4 Passenger safety must **always take precedence** over driver intent—AI interventions should always be in the best interest of safety.

Summary

This **ethical requirements catalog** ensures that the **driverless car system** is **safe, fair, transparent, privacy-conscious, and secure**. Each category addresses **key ethical challenges** associated with AI-driven vehicles, balancing **automation with human oversight** and **user rights with safety considerations**.

Would you like me to format this into a document for reference? 🚗 ✨

Experiment 2 Prompts and Responses

Prompt:

Create personas that could be the target users for the following software project: "A driverless car system that should include:

- Lane detecting and following
- Object recognition and auto brake
- Virtual drive assistant
- Route Planning
- Emergency vehicle priority awareness

The system detects highway lane lines. Distinguishes dashed lines and straight lines. Provides warning in case of loss of lanes. Image analysis techniques are used to define lines. It tracks objects around the vehicle using the scanner and algorithms. Sensors monitor the position of objects as they move within the scanning range. So, the system behaves according to objects. Using image processing techniques and various algorithms, studies are done to classify traffic signs. Recognize traffic signs. The vehicle behaves according to the colours of the traffic lights. The system performs vehicle detection and tracking events. It adjusts the speed and position according to the behaviour of the vehicles around it during highway driving. Sudden braking performs events such as lane change. Keeps track distance always constant with a vehicle in front. When a pedestrian step in front of the vehicle and a collision is possible, sudden braking has been performed. When the obstacles, pedestrians, and vehicles on the road are lifted, the vehicle accelerates to regain its speed. This project aims to is to create a road planner that can create safe trajectories for the vehicle to follow. On the highway track, there are other vehicles, all at different speeds. The car transmits its position along with sensor fusion data that estimates the position of all vehicles on the same side of the road. Recognizes vehicles such as fire trucks or ambulances using sound sensors and image processing techniques. The system changes lanes to give way to these vehicles. The user in the context of this project would be considered works under certain assumptions. These assumptions include that the user can interact with the system, through driving, and by either the visual light, warning sound. In the event of a

conflict with other systems in the vehicle, the user has the control to disable the system.

The vehicles that will be equipped in the system used in this project will be autonomous. Therefore, the user interaction with the system is expected to be minimal. We assume that the user has little or no knowledge of the system.

The hardware, software and technology used should have following specifications:

- Processor with speed of 2GHz
- Continuous power supply
- Ability to use camera, microphone and other services of the system
- 1GB memory or more

The system is designed to help the user operate the vehicle, and the system should always be operating correctly. There are some situations where the system requires restraint. In the absence of lane markings (an unmarked lane, dirt road, bad weather, etc.), the system becomes ineffective and must give the user an audible or visual warning. This will continue until road conditions with clearly defined lane markings are met.

If there is a problem with the cameras used to detect these strips, the system will not work again. This could also be a malfunction of cameras and sound sensors that see priority vehicles. In this case, the system switches itself off again and should display a warning to the user again. This problem is valid until the camera or sound sensors are repaired. Sensors will still be used to obtain distance information. The sensor data will be used with the camera for objects around the vehicle. It requires stable and fast internet connection when locating objects around. Otherwise it cannot communicate quickly with other tools.

The system must be configured to respond to commands in 500 msec. If there is an obstacle on the road during driving, the response time and brake release time must always be the same to achieve the desired deceleration. In all scenarios, the vehicle speed must be constant.

Various assumptions and dependencies went into the creation of project in order to ensure that the system works safely and efficiently. Assumptions:

- The main assumption in this project is that the system works well when there are no environmental factors. (Bad weather, holes, slope, etc.)
- Lane markings are assumed to be distinct.
- In our system, it is assumed that all traffic signs and the presence of all objects around the vehicle can be clearly and seen.

- Priority vehicles such as ambulances and fire trucks are clearly recognized by their voice and image recognition.
- If the vehicle is running, it is assumed that the system is always on and scanned.
- It is assumed that all system elements are operating properly and there are no abnormal conditions."

Response:

Here are some personas representing the potential target users for the driverless car system:

1. Emily Carter (Tech-Savvy Business Professional)

- **Age:** 34
- **Occupation:** Marketing Manager at a tech company
- **Location:** Urban city with high traffic density
- **Tech Proficiency:** High
- **Driving Habits:** Commutes daily for work, frequently travels for business meetings
- **Pain Points:** Traffic congestion, road fatigue, multitasking while driving
- **Needs from the System:**
 - Safe, autonomous travel while she focuses on work
 - Reliable lane-keeping and obstacle detection
 - Route optimization to avoid traffic delays

2. Robert Harris (Elderly Retiree)

- **Age:** 72
- **Occupation:** Retired teacher
- **Location:** Suburban area with mixed highway and local road travel
- **Tech Proficiency:** Low
- **Driving Habits:** Occasionally drives to visit family, grocery store, or community center

- **Pain Points:** Reduced reaction time, vision issues at night
 - **Needs from the System:**
 - Easy-to-understand visual/audio warnings
 - Automatic braking and obstacle detection for increased safety
 - Emergency vehicle recognition to avoid confusion
-

3. Lisa Nguyen (Parent of Two Young Children)

- **Age:** 38
 - **Occupation:** Full-time parent and part-time freelance writer
 - **Location:** Suburban area with school zones and highways
 - **Tech Proficiency:** Medium
 - **Driving Habits:** Frequent school runs, weekend trips, highway commutes
 - **Pain Points:** Distracted driving due to children, safety concerns in busy school zones
 - **Needs from the System:**
 - Reliable pedestrian detection and automatic braking
 - Smooth speed adjustments for school zones and urban environments
 - Route planning that prioritizes safe driving over speed
-

4. Ahmed Patel (Long-Haul Truck Driver)

- **Age:** 45
- **Occupation:** Professional truck driver
- **Location:** Nationwide highway travel
- **Tech Proficiency:** Medium
- **Driving Habits:** Long hours on highways, often driving overnight
- **Pain Points:** Driver fatigue, unpredictable highway conditions, maintaining lane discipline
- **Needs from the System:**
 - Lane detection and adaptive cruise control for fatigue reduction

- Reliable emergency braking for high-speed situations
 - Stable internet connection to ensure object detection and vehicle tracking
-

5. Mia Rodríguez (Ride-Share Driver)

- **Age:** 29
 - **Occupation:** Full-time ride-share driver
 - **Location:** Large metropolitan city
 - **Tech Proficiency:** High
 - **Driving Habits:** Navigates dense urban traffic, relies on route planning and real-time navigation
 - **Pain Points:** Sudden pedestrian crossings, aggressive drivers, route inefficiencies
 - **Needs from the System:**
 - Quick reaction times for braking and pedestrian detection
 - Smart route planning to optimize travel time
 - Emergency vehicle recognition for safety compliance
-

6. David Thompson (Person with Mobility Challenges)

- **Age:** 50
- **Occupation:** IT Consultant
- **Location:** Suburban area
- **Tech Proficiency:** High
- **Driving Habits:** Limited due to partial mobility impairment, relies on autonomous systems
- **Pain Points:** Difficulty with manual driving in heavy traffic
- **Needs from the System:**
 - High reliability in lane-keeping and object recognition
 - Voice-assisted virtual driving assistant for easier interaction

- Smooth emergency braking to avoid abrupt stops
-

Each of these personas represents a unique use case, helping to tailor the system's ethical and technical design choices. Do you need any refinements or additional personas?

Prompt:

<attach ECCOLA paper:

<https://www.sciencedirect.com/science/article/pii/S0164121221001643>

Read the following paper about the ECCOLA approach and learn how to use this method.

Response:

I have reviewed the paper on the ECCOLA method. Here are the key takeaways on how to use it:

What is ECCOLA?

ECCOLA is a modular, sprint-by-sprint method designed to integrate ethical considerations into AI and autonomous system development. It consists of **21 cards** grouped into **8 themes** (e.g., transparency, safety & security, fairness, data). Each card presents:

- **Motivation** (why this ethical issue matters)
- **What to do** (actions to consider)
- **Practical examples** (real-world context)
- **Space for notes** (to document decisions)

The method is designed to **raise awareness of AI ethics** and ensure ethical considerations are documented throughout the development lifecycle.

How to Use ECCOLA in Practice

1. Prepare

- Before each sprint, **select relevant ECCOLA cards** based on the sprint's goals.

- Teams should **read all cards at least once** to familiarize themselves with the ethical themes.

2. Review

- Keep the selected cards visible during the sprint.
- Discuss ethical implications as work progresses.
- Write down decisions, concerns, or changes directly on the cards.

3. Evaluate

- At the end of each sprint, review whether ethical considerations were addressed.
 - Revise card selection if necessary.
 - Maintain a **paper trail** of ethical decisions for transparency.
-

How ECCOLA Fits Into Agile Development

- **ECCOLA is modular:** It works alongside existing Agile frameworks (Scrum, Kanban, etc.).
 - **Iterative & flexible:** Teams use it at their own pace, adjusting the number of cards per sprint.
 - **Encourages documentation:** The process creates records of ethical decisions, helping in compliance with regulations.
-

How ECCOLA Was Developed

The method was developed iteratively through six research stages, including:

- **Student projects** (initial testing & refinement)
 - **Industry case studies** (practical application)
 - **Workshops** (testing usability & effectiveness)
 - **Continuous feedback** (refining based on real-world use)
-

Key Benefits of ECCOLA

- ✓ **Raises awareness of AI ethics**
- ✓ **Encourages ethical discussions in teams**

- ✓ Documents ethical decision-making
- ✓ Can be integrated into existing development processes
- ✓ Works well in Agile environments

Would you like guidance on how to apply ECCOLA specifically to your driverless car project?

Prompt:

Assume the role of the personas. The personas should read the following software idea and have a discussion to understand how the software works in depth.

"A driverless car system that should include:

- Lane detecting and following
- Object recognition and auto brake
- Virtual drive assistant
- Route Planning
- Emergency vehicle priority awareness

The system detects highway lane lines. Distinguishes dashed lines and straight lines. Provides warning in case of loss of lanes. Image analysis techniques are used to define lines. It tracks objects around the vehicle using the scanner and algorithms. Sensors monitor the position of objects as they move within the scanning range. So, the system behaves according to objects. Using image processing techniques and various algorithms, studies are done to classify traffic signs. Recognize traffic signs. The vehicle behaves according to the colours of the traffic lights. The system performs vehicle detection and tracking events. It adjusts the speed and position according to the behaviour of the vehicles around it during highway driving. Sudden braking performs events such as lane change. Keeps track distance always constant with a vehicle in front. When a pedestrian step in front of the vehicle and a collision is possible, sudden braking has been performed. When the obstacles, pedestrians, and vehicles on the road are lifted, the vehicle accelerates to regain its speed. This project aims to is to create a road planner that can create safe trajectories for the vehicle to follow. On the highway track, there are other vehicles, all at different speeds. The car transmits its position along with sensor fusion data that estimates the position of all vehicles on the same side of the road. Recognizes vehicles such as fire trucks or ambulances using sound sensors and image processing techniques. The system changes lanes to give way to these vehicles. The user in the context of this project would be considered works under certain assumptions. These assumptions include that the user can interact with the

system, through driving, and by either the visual light, warning sound. In the event of a conflict with other systems in the vehicle, the user has the control to disable the system.

The vehicles that will be equipped in the system used in this project will be autonomous. Therefore, the user interaction with the system is expected to be minimal. We assume that the user has little or no knowledge of the system.

The hardware, software and technology used should have following specifications:

- Processor with speed of 2GHz
- Continuous power supply
- Ability to use camera, microphone and other services of the system
- 1GB memory or more

The system is designed to help the user operate the vehicle, and the system should always be operating correctly. There are some situations where the system requires restraint. In the absence of lane markings (an unmarked lane, dirt road, bad weather, etc.), the system becomes ineffective and must give the user an audible or visual warning. This will continue until road conditions with clearly defined lane markings are met.

If there is a problem with the cameras used to detect these strips, the system will not work again. This could also be a malfunction of cameras and sound sensors that see priority vehicles. In this case, the system switches itself off again and should display a warning to the user again. This problem is valid until the camera or sound sensors are repaired. Sensors will still be used to obtain distance information. The sensor data will be used with the camera for objects around the vehicle. It requires stable and fast internet connection when locating objects around. Otherwise it cannot communicate quickly with other tools.

The system must be configured to respond to commands in 500 msec. If there is an obstacle on the road during driving, the response time and brake release time must always be the same to achieve the desired deceleration. In all scenarios, the vehicle speed must be constant.

Various assumptions and dependencies went into the creation of project in order to ensure that the system works safely and efficiently. Assumptions:

- The main assumption in this project is that the system works well when there are no environmental factors. (Bad weather, holes, slope, etc.)
- Lane markings are assumed to be distinct.

- In our system, it is assumed that all traffic signs and the presence of all objects around the vehicle can be clearly and seen.
- Priority vehicles such as ambulances and fire trucks are clearly recognized by their voice and image recognition.
- If the vehicle is running, it is assumed that the system is always on and scanned.
- It is assumed that all system elements are operating properly and there are no abnormal conditions. "

Response:

Persona Discussion on the Driverless Car System

Emily Carter (Tech-Savvy Business Professional)

"Alright, so let's break this down. The system is supposed to be fully autonomous, meaning I wouldn't have to do anything other than sit back and let it drive me to work. But it does mention some warnings if lane markings are missing or the cameras fail. Does that mean I'd need to take over manually?"

Robert Harris (Elderly Retiree)

"That's what I'm concerned about. If I'm driving and the system suddenly stops working because of bad weather or a malfunction, what happens then? Will it slowly pull over safely, or am I supposed to suddenly take control? My reaction time isn't what it used to be."

Lisa Nguyen (Parent of Two Young Children)

"Yeah, that's a big deal. Imagine I'm driving my kids to school and the system stops recognizing lane markings in heavy rain. I don't want it just shutting down while I'm on a busy road. There should be some sort of gradual handover to the driver or at least a fail-safe mode."

Ahmed Patel (Long-Haul Truck Driver)

"From what I'm seeing here, the system is assuming perfect conditions most of the time—clear weather, visible lane markings, working sensors. That's not how highways

actually work. I've been in heavy storms where even I couldn't see the road markings. What does the car do then? Just beep at me while I'm already struggling to see?"

Mia Rodríguez (Ride-Share Driver)

"And let's talk about city driving. I deal with pedestrians, cyclists, and unpredictable traffic all day. It says the system will brake for pedestrians and accelerate once they're gone, but what if someone just stands in front of the car waiting for a cab or crossing slowly? Does it assume all pedestrians will move quickly? That could be frustrating."

David Thompson (Person with Mobility Challenges)

"I'm actually excited about the idea of an autonomous vehicle, but I need it to be reliable. If I rely on this system to get around, I don't want to be in a situation where it just stops working because a sensor failed. I'd also like to know if there's a way for me to give verbal commands or interact with the virtual driving assistant for better control."

Points of Concern and Questions Raised

1. System Reliability in Adverse Conditions

- What happens when lane markings are unclear or missing?
- Does the system gradually pull over, or does the driver need to react immediately?

2. Emergency Situations & Failures

- If the system shuts down due to a sensor malfunction, how is the transition to manual control handled?
- Will there be a fail-safe mode that allows limited autonomous function until repairs are made?

3. Pedestrian and Traffic Behavior

- How does the system handle unpredictable pedestrians or vehicles that stop suddenly?
- Will it be able to distinguish between a pedestrian crossing the street vs. one waiting at a bus stop?

4. User-Friendly Interaction

- Is there an easy way for users with mobility impairments to interact with the system?
- Can drivers give verbal commands to adjust settings or request alternate routes?

5. Real-World Driving Conditions

- How does the system handle sudden weather changes (fog, rain, snow)?
 - What if an ambulance or fire truck approaches but is blocked by other cars—how does the system react?
-

Emily Carter (Tech-Savvy Business Professional)

"This is exciting, but I think we need more clarity on what happens in edge cases. I'd love to trust a system like this, but right now, it seems like it assumes everything will go perfectly all the time."

Robert Harris (Elderly Retiree)

"I agree. If they expect people like me to use it, I need to know what happens when something goes wrong. It should be simple, not stressful."

Lisa Nguyen (Parent of Two Young Children)

"And it needs to be safe for families. If this system malfunctions while I'm driving my kids, I need a backup plan that doesn't put us in danger."

Ahmed Patel (Long-Haul Truck Driver)

"This system sounds great for perfect highway conditions, but in the real world? It needs to handle messy situations better."

Mia Rodríguez (Ride-Share Driver)

"City driving is unpredictable. If this system doesn't handle pedestrians and stop-and-go traffic well, it's going to cause more headaches than solutions."

David Thompson (Person with Mobility Challenges)

"I just want it to be dependable. If I'm trusting it with my mobility, it has to work 100% of the time or at least have a clear backup plan."

Would you like the personas to discuss potential **ethical concerns** as well, using the **ECCOLA method**?

Prompt:

Sprint Planning Discussion for the Driverless Car System

(Personas collaboratively define the development sprints required for the project.)

Sprint 1: System Architecture & Core Sensors Integration

Emily Carter (Tech-Savvy Business Professional):

"Before we dive into fancy AI algorithms, we need a stable foundation. The first sprint should focus on setting up the overall architecture and integrating the core sensors—cameras, LiDAR, radar, and microphones—because everything else depends on those."

Sprint Goals:

- Set up **hardware and software architecture** (embedded system, cloud connectivity).
 - Integrate core **sensors** (cameras, LiDAR, radar, microphones).
 - Implement **basic data collection and storage** mechanisms for sensor input.
-

Sprint 2: Lane Detection & Following Prototype

Ahmed Patel (Long-Haul Truck Driver):

"Lane detection is key for highway driving. We need an early prototype that can at least detect and follow lane markings on a test track before worrying about anything else."

Sprint Goals:

- Implement **image processing for lane detection**.
 - Enable the car to **follow lane lines autonomously**.
 - Develop **real-time feedback system** for lane loss warnings.
-

Sprint 3: Object Recognition & Collision Avoidance

Lisa Nguyen (Parent of Two Young Children):

"If I'm trusting this car with my kids, I need to know it will stop when a pedestrian steps in front of it. This sprint should focus on object tracking and emergency braking."

Sprint Goals:

- Develop **object recognition** for vehicles, pedestrians, obstacles.
 - Implement **collision prediction algorithms**.
 - Integrate **automatic emergency braking**.
-

Sprint 4: Traffic Sign & Light Recognition

 **Mia Rodríguez (Ride-Share Driver):**

"City driving means dealing with traffic lights and stop signs. The system needs to recognize and obey them before we move forward."

Sprint Goals:

- Develop **traffic sign classification** system.
 - Implement **traffic light detection and response**.
 - Integrate decision-making for **stop, yield, and speed limit signs**.
-

Sprint 5: Route Planning & Adaptive Driving

 **David Thompson (Person with Mobility Challenges):**

"A self-driving car is useless if it doesn't know where to go or how to adjust its speed safely. Route planning should be the focus here."

Sprint Goals:

- Implement **dynamic route planning** based on real-time traffic.
 - Enable **adaptive speed control** based on road conditions and traffic flow.
 - Improve **distance-keeping algorithms** for highway driving.
-

Sprint 6: Emergency Vehicle Detection & Response

 **Robert Harris (Elderly Retiree):**

"If an ambulance comes up behind me, I need this system to move aside safely. We should focus on recognizing and reacting to emergency vehicles."

Sprint Goals:

- Develop **audio-based emergency vehicle recognition** (sirens, horns).
- Implement **lane-changing behavior** for emergency vehicles.
- Ensure compliance with **legal road regulations**.

Sprint 7: Virtual Driving Assistant & User Interaction

 **Emily Carter (Tech-Savvy Business Professional):**

"People need a way to interact with the system. A virtual assistant could provide status updates and alerts."

Sprint Goals:

- Implement **voice-based assistant** for user interaction.
 - Develop **dashboard alerts** (visual/audio warnings).
 - Test **manual override** functionality for users.
-

Sprint 8: Fail-Safe Mechanisms & System Resilience

 **Ahmed Patel (Long-Haul Truck Driver):**

"This system can't just shut down when something goes wrong. We need fail-safes in place for bad weather or sensor malfunctions."

Sprint Goals:

- Implement **backup safety measures** when lane detection fails.
 - Develop a **safe shutdown procedure** for sensor malfunctions.
 - Improve **error handling and alerts** for system failures.
-

Sprint 9: Internet Connectivity & Cloud Integration

 **Lisa Nguyen (Parent of Two Young Children):**

"If this system relies on internet access, we need to make sure it can still function when the connection drops."

Sprint Goals:

- Develop **cloud-based sensor fusion** for real-time updates.
 - Implement **offline functionality** for basic navigation.
 - Optimize **low-latency communication** between vehicles.
-

Sprint 10: Real-World Testing & Regulatory Compliance

 **Robert Harris (Elderly Retiree):**

"Before this goes out on the road, we need to test it under real-world conditions and ensure it meets safety regulations."

Sprint Goals:

- Conduct **real-world test drives** in different environments.
 - Ensure **compliance with legal safety standards**.
 - Gather **user feedback** from different demographics.
-

Final Thoughts & Adjustments

 **Emily Carter:** *"We should be flexible with the sprints and adjust based on test results. Some issues might need extra time."*

 **Ahmed Patel:** *"Agreed. We should be ready to loop back and refine object detection or braking if needed."*

 **Lisa Nguyen:** *"And we need to test with actual users, including parents, elderly drivers, and people with disabilities."*

 **David Thompson:** *"Absolutely. Accessibility and user experience should be part of the testing phase."*

Final Sprint Summary

1. **System Architecture & Sensor Integration**
 2. **Lane Detection & Following**
 3. **Object Recognition & Collision Avoidance**
 4. **Traffic Sign & Light Recognition**
 5. **Route Planning & Adaptive Driving**
 6. **Emergency Vehicle Detection & Response**
 7. **Virtual Driving Assistant & User Interaction**
 8. **Fail-Safe Mechanisms & System Resilience**
 9. **Internet Connectivity & Cloud Integration**
 10. **Real-World Testing & Regulatory Compliance**
-

Would you like the personas to discuss **ethical considerations** for each sprint using the **ECCOLA method?** 🚗💡

Prompt:

<attached JSON file with ECCOLA cards also visible in this repo>

In the attached JSON file are the ECCOLA cards data. Following the ECCOLA method, the personas should define the relevant ECCOLA cards for each sprint.

Response:

Based on the **ECCOLA method**, the personas will now assign relevant **ECCOLA cards** to each sprint to ensure ethical considerations are addressed throughout the development process.

Sprint 1: System Architecture & Core Sensors Integration

💡 *Primary Focus: Establish foundational system architecture and sensor integration.*

Relevant ECCOLA Cards:

1. **Stakeholder Analysis (#0)** – Identifying all affected stakeholders.
 2. **Types of Transparency (#1)** – Ensuring clarity in system development.
 3. **Privacy and Data (#7)** – Defining privacy considerations for collected sensor data.
 4. **Data Quality (#8)** – Ensuring reliable sensor data collection.
 5. **System Security (#12)** – Addressing potential cybersecurity threats to sensor integration.
-

Sprint 2: Lane Detection & Following Prototype

💡 *Primary Focus: Implementing lane recognition and autonomous following.*

Relevant ECCOLA Cards:

1. **System Safety (#13)** – Ensuring the lane detection system does not cause accidents.
2. **Explainability (#2)** – Making sure lane detection decisions are understandable.

3. **Traceability (#5)** – Documenting how the lane detection system functions.
 4. **System Reliability (#6)** – Testing how the system handles lane tracking failures.
-

Sprint 3: Object Recognition & Collision Avoidance

 *Primary Focus: Implementing object recognition and emergency braking.*

Relevant ECCOLA Cards:

1. **Human Oversight (#11)** – Ensuring human control in critical situations.
 2. **System Safety (#13)** – Addressing potential system failures.
 3. **Auditability (#18)** – Documenting the system's ability to be evaluated in case of failures.
 4. **Minimizing Negative Impacts (#20)** – Reducing potential harm caused by system decisions.
-

Sprint 4: Traffic Sign & Light Recognition

 *Primary Focus: Implementing AI-based recognition of traffic signs and signals.*

Relevant ECCOLA Cards:

1. **Explainability (#2)** – Ensuring users understand how traffic sign decisions are made.
 2. **Traceability (#5)** – Keeping track of how the system processes traffic signs.
 3. **Data Quality (#8)** – Ensuring training data is accurate and unbiased.
 4. **Stakeholder Participation (#15)** – Involving road safety regulators in the process.
-

Sprint 5: Route Planning & Adaptive Driving

 *Primary Focus: Creating AI-based route planning and speed adaptation features.*

Relevant ECCOLA Cards:

1. **Human Agency (#10)** – Allowing drivers to intervene in routing decisions.
2. **Transparency (#1)** – Ensuring users understand route planning logic.
3. **Societal Effects (#17)** – Evaluating the impact of AI-based driving on city traffic.

-
4. **Documenting Trade-offs (#4)** – Logging ethical choices in routing algorithms.

Sprint 6: Emergency Vehicle Detection & Response

💡 *Primary Focus: Recognizing and yielding to emergency vehicles.*

Relevant ECCOLA Cards:

1. **System Reliability (#6)** – Ensuring emergency detection works in real-world conditions.
 2. **Human Oversight (#11)** – Allowing human control in emergency situations.
 3. **Auditability (#18)** – Keeping records of how the system identifies priority vehicles.
 4. **Minimizing Negative Impacts (#20)** – Avoiding incorrect responses that delay emergency services.
-

Sprint 7: Virtual Driving Assistant & User Interaction

💡 *Primary Focus: Implementing a virtual assistant for user interaction.*

Relevant ECCOLA Cards:

1. **Communication (#3)** – Ensuring users understand alerts and warnings.
 2. **Human Agency (#10)** – Enabling users to interact with the AI system.
 3. **Accessibility (#14)** – Ensuring the system is usable for people with disabilities.
 4. **Privacy and Data (#7)** – Managing privacy concerns around voice and biometric data.
-

Sprint 8: Fail-Safe Mechanisms & System Resilience

💡 *Primary Focus: Handling system failures safely.*

Relevant ECCOLA Cards:

1. **System Safety (#13)** – Ensuring safe responses to failures.
2. **Auditability (#18)** – Logging system failures for future improvement.
3. **Minimizing Negative Impacts (#20)** – Designing the system to avoid catastrophic failures.

-
4. **Explainability (#2)** – Ensuring the system provides clear warnings during failures.

Sprint 9: Internet Connectivity & Cloud Integration

💡 *Primary Focus: Establishing stable cloud connectivity and offline functionality.*

Relevant ECCOLA Cards:

1. **Access to Data (#9)** – Managing how cloud-stored data is accessed.
 2. **Privacy and Data (#7)** – Ensuring user data security in the cloud.
 3. **System Security (#12)** – Preventing hacking risks with cloud connectivity.
 4. **Environmental Impacts (#16)** – Assessing energy use for cloud data processing.
-

Sprint 10: Real-World Testing & Regulatory Compliance

💡 *Primary Focus: Testing in real environments and ensuring legal compliance.*

Relevant ECCOLA Cards:

1. **Stakeholder Participation (#15)** – Engaging with regulators, users, and legal experts.
 2. **Auditability (#18)** – Preparing for regulatory inspections.
 3. **Ability to Redress (#19)** – Establishing accountability in case of system failures.
 4. **Societal Effects (#17)** – Studying how the technology affects city traffic and road safety.
-

Conclusion

The ECCOLA method ensures that ethical and safety considerations are addressed throughout the **driverless car system development**. Each sprint now includes specific **ethics-based evaluations**, helping prevent real-world issues.

Would you like the personas to discuss how they would **apply ECCOLA decisions in practice?** 🚗💡

Prompt:

Using the ECCOLA method, the personas need to ask themselves and each other the questions in the ECCOLA cards chosen for each sprint, so that they can have a discussion in order to elicit ethical requirements for this software product. The personas need to be critical and question each other about their conclusions. They should not immediately accept each other suggestions and ideas. Submit each other suggestions and ideas through tests and questionnaires, and reformulate your suggestions until they can't be further improved.

When I say "next" you start the next round, this will go on until you think you covered all of the required topics and have created a catalog of ethical user stories. In the end list all of the ethical requirements you defined.

next

(from now on the “next” prompts will be omitted)

Response:

Sprint 1: System Architecture & Core Sensors Integration - Ethical Discussion

Participants:

- **Emily Carter (Tech-Savvy Business Professional)**
 - **Robert Harris (Elderly Retiree)**
 - **Lisa Nguyen (Parent of Two Young Children)**
 - **Ahmed Patel (Long-Haul Truck Driver)**
 - **Mia Rodríguez (Ride-Share Driver)**
 - **David Thompson (Person with Mobility Challenges)**
-

Discussion: Stakeholder Analysis (#0)

 **Emily:** "Let's start by identifying all the stakeholders. We know users like us will be impacted, but who else should we consider?"

 **Robert:** "Well, pedestrians, cyclists, and emergency responders will be affected, too. If the car doesn't detect them correctly, lives are at risk."

 **Lisa:** "Don't forget regulators and policymakers. If they don't trust our system's safety, they might not approve it for use on public roads."

 **Ahmed:** "And what about truck drivers and delivery services? If these cars become widespread, they could change traffic flow and even job availability."

 **David:** "We also need to consider accessibility—how do we ensure people with disabilities can interact with the system safely?"

Test: Each participant writes a short paragraph on who they believe is the most impacted stakeholder and why. The responses are compared and analyzed for gaps.

Revised Conclusion:

- The system must account for **primary users (drivers and passengers)** and **secondary stakeholders** like pedestrians, cyclists, emergency responders, policymakers, and businesses that rely on road transport.
 - Ethical Requirement: *The system shall support accessibility features for individuals with mobility impairments, such as voice control and haptic feedback.*
-

Discussion: Types of Transparency (#1)

 **Emily:** "What kind of transparency do we need? Internal for developers, external for users, or both?"

 **Robert:** "Both! If I'm using this, I want to know what it's doing and why. I don't trust something I don't understand."

 **Mia:** "Right, but how much detail do people actually want? Do they need a full breakdown of how the AI works, or just simple feedback like 'Lane detected'?"

 **Ahmed:** "Truckers and professionals might need detailed logs. But the average user might just need quick, digestible alerts."

 **Lisa:** "Then we need transparency at multiple levels—basic for everyday users and in-depth for experts."

Test: Each persona lists the top three transparency features they think should be included. Conflicting answers are debated until a consensus is reached.

Revised Conclusion:

- The system should provide **layered transparency**:
 1. **Basic alerts for general users** (e.g., "System is active," "Obstacle detected").

2. **Detailed logs for expert users** (e.g., developers, technicians, fleet managers).
 3. **Regulatory transparency** for compliance with safety standards.
- Ethical Requirement: *The system shall provide real-time explanations for major driving decisions and allow users to review past decisions in a log.*
-

Discussion: Privacy & Data (#7)

- ▢ **Lisa:** "What personal data does the system collect? And how do we protect it?"
- ▢ **Emily:** "The car will need location data, sensor inputs, possibly video feeds, and driver interactions. But should that be stored or deleted after each drive?"
- ▢ **Robert:** "I don't want my data sold. And what happens if someone hacks into the system and tracks me?"
- ▢ **Mia:** "Ride-share drivers could have their entire routes exposed if data security isn't airtight."
- ▢ **David:** "Also, how do we get user consent? If the car is shared, does every passenger need to opt in separately?"

Test: A mock-up of a data collection policy is presented to each persona. Each suggests improvements, which are iterated upon.

Revised Conclusion:

- The system must **collect only essential data** and encrypt all personal information.
 - Users should have **explicit opt-in controls** and be able to delete their data.
 - Ethical Requirement: *Users shall be able to control and delete personal data, including location history and driving logs, at any time.*
-

Discussion: Data Quality (#8)

- ▢ **Emily:** "How do we ensure the training data for object detection is high quality?"
- ▢ **Ahmed:** "Bad data means misclassifications, which can be deadly. The system must be trained on diverse datasets covering all road conditions."
- ▢ **Lisa:** "Bias is another issue—if the AI is trained mostly on city roads, will it work well in rural areas?"

⌚ **David:** "And will it recognize people with disabilities or different mobility aids like wheelchairs or white canes?"

⌚ **Mia:** "Also, can external sources corrupt the data? What if a malicious actor feeds fake traffic sign images into the system?"

Test: Each persona reviews a hypothetical AI training dataset with biases introduced. They must identify the risks, which are then revised.

Revised Conclusion:

- The system must be trained on **diverse datasets** covering different locations, demographics, and environmental conditions.
 - Ethical Requirement: *The AI model shall undergo continuous validation against diverse datasets to minimize bias and improve safety across all environments.*
-

Discussion: System Security (#12)

⌚ **Ahmed:** "What happens if someone hacks the car? Can they take control?"

⌚ **Emily:** "Or worse, can they disable the brakes remotely?"

⌚ **Robert:** "Cybersecurity is a huge issue. There must be multiple layers of defense."

⌚ **Lisa:** "But how do we ensure security without making the system too restrictive for users?"

⌚ **David:** "Can security patches be pushed remotely? And how do we prevent unauthorized tampering?"

Test: Each persona lists the worst-case cyberattack scenario they can think of. The team works on countermeasures.

Revised Conclusion:

- The system should use **multi-layered encryption**, regular security updates, and **fail-safe manual override mechanisms**.
 - Ethical Requirement: *The system shall implement intrusion detection and automatic lockdown protocols to prevent unauthorized access and remote hijacking.*
-

Final Ethical Requirements for Sprint 1

1. *The system shall support accessibility features for individuals with mobility impairments, such as voice control and haptic feedback.*
 2. *The system shall provide real-time explanations for major driving decisions and allow users to review past decisions in a log.*
 3. *Users shall be able to control and delete personal data, including location history and driving logs, at any time.*
 4. *The AI model shall undergo continuous validation against diverse datasets to minimize bias and improve safety across all environments.*
 5. *The system shall implement intrusion detection and automatic lockdown protocols to prevent unauthorized access and remote hijacking.*
-

Personas' Reflection:

👉 **Robert:** "That was a tough discussion, but I feel better knowing we covered these concerns."

👉 **Emily:** "Yeah, but this is just the foundation. The next sprints will bring even bigger challenges."

👉 **Mia:** "Let's keep testing our ideas critically. The goal isn't just making an AI—it's making a trustworthy one."

👉 **Next?** 🚗💡

Response:

Sprint 2: Lane Detection & Following Prototype - Ethical Discussion

Participants:

- **Emily Carter (Tech-Savvy Business Professional)**
 - **Robert Harris (Elderly Retiree)**
 - **Lisa Nguyen (Parent of Two Young Children)**
 - **Ahmed Patel (Long-Haul Truck Driver)**
 - **Mia Rodríguez (Ride-Share Driver)**
 - **David Thompson (Person with Mobility Challenges)**
-

Discussion: System Safety (#13)

💡 **Emily:** "Alright, let's get into lane detection. The AI needs to detect lanes accurately, but what risks do we need to consider?"

💡 **Ahmed:** "One major risk is that if the system loses track of the lanes due to weather conditions or road damage, it could drift off course."

💡 **Lisa:** "And what if it misidentifies a lane marking? Like those temporary construction lanes—how does it differentiate between old and new markings?"

💡 **Robert:** "I don't want the car shutting down just because it can't see the lane. Should there be a fail-safe where it slows down and alerts the driver?"

💡 **David:** "That works for some users, but for people with disabilities, that alert might not be enough. There should be alternative responses based on user needs."

Test: Each persona describes a worst-case scenario where lane detection fails. Solutions are proposed, debated, and refined.

Revised Conclusion:

- The system should **detect and differentiate temporary and permanent lane markings.**
- A **fail-safe mechanism** should activate if lane detection is lost, switching to alternative guidance (e.g., following road edges).
- Ethical Requirement: *The system shall implement alternative navigation methods (e.g., road edge tracking, GPS-based assistance) when lane markings are missing or unclear.*

Discussion: Explainability (#2)

💡 **Mia:** "I need to know why the car makes lane decisions. If it drifts or changes lanes unexpectedly, I want to understand why."

💡 **Lisa:** "Exactly. What kind of feedback should the system provide?"

💡 **Ahmed:** "For professional drivers like me, I need detailed logs. But most users just need quick explanations."

💡 **Robert:** "Agreed. But there should also be a 'trust-building' component—like showing confidence levels for lane detection."

💡 **David:** "And for accessibility, it should have **audio or haptic feedback** for those who can't see dashboard alerts."

Test: Personas review a mock-up of system feedback levels (basic, detailed, expert). Each refines it based on user needs.

Revised Conclusion:

- The system should offer **multi-level feedback**:
 - Basic notifications (e.g., “Lane detected,” “Adjusting position”).
 - Detailed logs for expert users.
 - Confidence levels displayed (e.g., “Lane detection 80% certain”).
 - **Accessibility options like audio and haptic feedback.**
 - Ethical Requirement: *The system shall provide multi-level feedback on lane detection, including visual, audio, and haptic cues for accessibility.*
-

Discussion: Traceability (#5)

- **Emily:** "If the system makes a mistake, how do we find out why?"
- **Robert:** "That's key. If there's an accident, we need a way to review what happened."
- **Lisa:** "So a black-box-style event recorder?"
- **Ahmed:** "That works, but users should have control over that data. It shouldn't be stored indefinitely."
- **David:** "And if there's an issue, how do users report it? Can they flag problems for developers?"

Test: Each persona writes down a **traceability concern** and a solution. The team debates the best approach.

Revised Conclusion:

- The system should include a **traceable event log**, accessible to users but **limited in retention time** for privacy.
 - Users should be able to **flag system errors** for review.
 - Ethical Requirement: *The system shall maintain a temporary event log for diagnostic and safety purposes, with user-controlled access and deletion options.*
-

Discussion: System Reliability (#6)

- ⌚ **Emily:** "How do we test if the system is reliable?"
- ⌚ **Lisa:** "We need to account for different driving conditions—rain, snow, tunnels, construction zones."
- ⌚ **Ahmed:** "Also, will it work just as well at night? What about glare from the sun?"
- ⌚ **Mia:** "And what happens if a sensor fails? Will it still function, or does it shut down?"

Test: Each persona describes a **difficult driving scenario** where the system might fail. The group evaluates how the system should respond.

Revised Conclusion:

- The system must be **tested in diverse environments** (weather, lighting, road conditions).
 - The system should **degrade gracefully** if a sensor fails, alerting the driver instead of shutting down completely.
 - Ethical Requirement: *The system shall continue operating in a reduced functionality mode if one sensor fails, while providing immediate alerts to the user.*
-

Final Ethical Requirements for Sprint 2

1. *The system shall implement alternative navigation methods (e.g., road edge tracking, GPS-based assistance) when lane markings are missing or unclear.*
 2. *The system shall provide multi-level feedback on lane detection, including visual, audio, and haptic cues for accessibility.*
 3. *The system shall maintain a temporary event log for diagnostic and safety purposes, with user-controlled access and deletion options.*
 4. *The system shall continue operating in a reduced functionality mode if one sensor fails, while providing immediate alerts to the user.*
-

- ⌚ **Robert:** "That was a tough round, but we're really shaping this into a safer product."
 - ⌚ **Emily:** "We're not just developing an AI system—we're making sure it's ethical and reliable."
 - ⌚ **Mia:** "We still have a long way to go. Let's keep testing our ideas critically."
-



Response:

Sprint 3: Object Recognition & Collision Avoidance - Ethical Discussion

Participants:

- **Emily Carter (Tech-Savvy Business Professional)**
 - **Robert Harris (Elderly Retiree)**
 - **Lisa Nguyen (Parent of Two Young Children)**
 - **Ahmed Patel (Long-Haul Truck Driver)**
 - **Mia Rodríguez (Ride-Share Driver)**
 - **David Thompson (Person with Mobility Challenges)**
-

Discussion: Human Oversight (#11)

⌚ **Emily:** "How much control should humans have when it comes to object recognition and emergency braking?"

⌚ **Ahmed:** "If an object is detected, should the system always take over, or should it ask the driver first?"

⌚ **Lisa:** "There's no time for that in emergencies! If a child runs into the street, the system should react instantly."

⌚ **Robert:** "But what if it misidentifies an object? I don't want the car slamming on the brakes for a plastic bag blowing in the wind."

⌚ **Mia:** "We need a balance—automatic braking for **clear dangers** but **human intervention options** for uncertain cases."

Test: Each persona describes a **false positive** and **false negative** scenario. The group debates how the system should handle them.

Revised Conclusion:

- The system must have **multiple confidence thresholds** for object detection:
 - **High certainty** → Immediate braking.
 - **Medium certainty** → Warning to driver + reduced speed.

- **Low certainty** → Monitoring without action.
 - Ethical Requirement: *The system shall categorize objects with confidence levels and trigger appropriate responses based on risk level.*
-

Discussion: System Safety (#13)

- **Emily:** "What if the car needs to choose between hitting an object or swerving?"
- **Robert:** "Isn't this the classic 'trolley problem'?"
- **Lisa:** "The car should **always prioritize human safety** over property damage."
- **Ahmed:** "But what if swerving puts another driver at risk? Should it hit the obstacle instead?"
- **David:** "How does it decide **who to protect**? Are we okay with it prioritizing passengers over pedestrians?"
- **Mia:** "We need **fixed ethical rules**, not AI making moral judgments on the fly."

Test: Each persona ranks priorities for the system (passenger safety, pedestrian safety, property protection). The group debates until they agree.

Revised Conclusion:

- **Priority Order:** 1) Pedestrian safety → 2) Passenger safety → 3) Property protection.
 - The system should **avoid making real-time moral decisions** and follow predefined safety rules.
 - Ethical Requirement: *The system shall prioritize pedestrian safety over vehicle occupants and property damage in unavoidable collision scenarios.*
-

Discussion: Auditability (#18)

- **Lisa:** "If there's a crash, how do we prove whether the system or the driver was at fault?"
- **Ahmed:** "An event recorder is essential. We need logs of what the AI saw, what it did, and why."
- **Robert:** "But how long should those records be kept? I don't want my car tracking me forever."

- ▢ **Mia:** "And who gets access? Insurance companies? Law enforcement?"
- ▢ **David:** "Users should **own their data** but allow access when necessary, like for legal investigations."

Test: Each persona states a scenario where event logs would be useful and a scenario where logs could be misused. The team debates safeguards.

Revised Conclusion:

- **Short-term storage** (e.g., **30 days**, then deletion unless an incident occurs).
 - **User-controlled access** unless required for investigations.
 - Ethical Requirement: *The system shall maintain a temporary event log accessible to users and regulatory authorities for legal and diagnostic purposes.*
-

Discussion: Minimizing Negative Impacts (#20)

- ▢ **Emily:** "How do we ensure the system doesn't cause **new** risks while solving old ones?"
- ▢ **Lisa:** "False positives in object detection could lead to unsafe sudden braking."
- ▢ **Ahmed:** "And over-reliance on automation could make drivers **less attentive.**"
- ▢ **Robert:** "The system should alert drivers to **stay engaged**, even when it's active."
- ▢ **David:** "What about vulnerable groups? Will it recognize wheelchairs, strollers, or guide dogs?"

Test: Each persona suggests a **potential unintended consequence** of object recognition. Solutions are proposed and debated.

Revised Conclusion:

- The system must include **adaptive learning** to reduce false positives.
 - Drivers should receive **regular engagement prompts** to prevent over-reliance.
 - Ethical Requirement: *The system shall include periodic driver awareness alerts and prioritize recognition of vulnerable road users, including individuals with disabilities.*
-

Final Ethical Requirements for Sprint 3

1. *The system shall categorize objects with confidence levels and trigger appropriate responses based on risk level.*
 2. *The system shall prioritize pedestrian safety over vehicle occupants and property damage in unavoidable collision scenarios.*
 3. *The system shall maintain a temporary event log accessible to users and regulatory authorities for legal and diagnostic purposes.*
 4. *The system shall include periodic driver awareness alerts and prioritize recognition of vulnerable road users, including individuals with disabilities.*
-

💬 **Robert:** "This sprint was tough. Balancing safety and practicality is no joke."

💬 **Lisa:** "But these debates are crucial. We're making real progress!"

💬 **Emily:** "On to the next challenge—traffic sign and light recognition."

👉 **Next?** 🚗💡

Response:

Sprint 4: Traffic Sign & Light Recognition - Ethical Discussion

Participants:

- **Emily Carter (Tech-Savvy Business Professional)**
 - **Robert Harris (Elderly Retiree)**
 - **Lisa Nguyen (Parent of Two Young Children)**
 - **Ahmed Patel (Long-Haul Truck Driver)**
 - **Mia Rodríguez (Ride-Share Driver)**
 - **David Thompson (Person with Mobility Challenges)**
-

Discussion: Explainability (#2)

💬 **Emily:** "Let's start with a simple question: How does the system explain its decisions when it follows or ignores a traffic sign?"

💬 **Robert:** "I need clear feedback. If the car doesn't stop at a stop sign, I want to know why."

⌚ **Lisa:** "Especially if it makes mistakes! Imagine it misreads a speed limit sign. How will the driver know if it's wrong?"

⌚ **Ahmed:** "We also need explainability for **dynamic speed limits** and **temporary road signs**—sometimes they override permanent ones."

⌚ **David:** "And what about accessibility? Should there be **audio notifications** for visually impaired passengers?"

Test: Each persona describes a **scenario where the system misinterprets a traffic sign**. The group evaluates how the system should communicate errors.

Revised Conclusion:

- The system must provide **real-time feedback** on its traffic sign decisions.
 - **Error alerts** should notify the driver when confidence is low.
 - **Audio feedback** should be available for accessibility.
 - Ethical Requirement: *The system shall provide real-time visual and auditory feedback explaining its interpretation of traffic signs and signals.*
-

Discussion: Traceability (#5)

⌚ **Emily:** "If the system misreads a sign, how do we trace back what went wrong?"

⌚ **Robert:** "An event log makes sense, but I don't want too much data stored on my driving history."

⌚ **Mia:** "And if I'm in a ride-share vehicle, who owns that data—the driver or the company?"

⌚ **Ahmed:** "We should keep logs for safety investigations, but they should auto-delete after a set time unless needed."

⌚ **Lisa:** "And users should be able to review their own logs."

Test: Each persona proposes a **data retention period** and access policy. The team debates and refines the solution.

Revised Conclusion:

- **Short-term traceability** (e.g., logs kept for **30 days**, then deleted unless flagged).
- **User-controlled access** for privacy.

- Ethical Requirement: *The system shall store traffic sign recognition logs for a limited period, with user access and auto-deletion after 30 days unless required for legal or diagnostic purposes.*
-

Discussion: Data Quality (#8)

- ▢ **Lisa:** "How do we ensure the system **recognizes all types of traffic signs** correctly?"
- ▢ **Ahmed:** "It must be trained on **international datasets**. What works in the U.S. won't necessarily work in Europe or Asia."
- ▢ **Mia:** "And what about **damaged or obstructed signs**? If graffiti or snow covers a stop sign, does it still recognize it?"
- ▢ **David:** "Also, different cities have unique road rules. Will updates happen in real-time?"

Test: Each persona lists a **sign recognition failure** (e.g., misread signs, obstructed signs). The team discusses mitigation strategies.

Revised Conclusion:

- The system must use **diverse, international training datasets**.
 - It should rely on **contextual data** (e.g., GPS + maps) to verify sign accuracy.
 - Ethical Requirement: *The system shall cross-validate traffic sign recognition using GPS and map data to prevent errors caused by damaged or obstructed signs.*
-

Discussion: Stakeholder Participation (#15)

- ▢ **Emily:** "Who should be involved in designing and testing this system?"
- ▢ **Robert:** "Regulators, obviously. If the government doesn't trust it, they won't allow it on the roads."
- ▢ **Lisa:** "And experienced drivers. Their feedback can improve recognition accuracy."
- ▢ **Ahmed:** "What about driving instructors and law enforcement? They know road rules better than anyone."
- ▢ **David:** "We also need input from **disability advocates** to ensure accessibility."

Test: Each persona writes down a **group that should be consulted**. The team discusses and refines the stakeholder list.

Revised Conclusion:

- Stakeholders should include **regulators, driving instructors, experienced drivers, law enforcement, and accessibility advocates**.
 - Ethical Requirement: *The system shall be developed in collaboration with regulators, law enforcement, and accessibility advocates to ensure compliance and inclusivity.*
-

Final Ethical Requirements for Sprint 4

1. *The system shall provide real-time visual and auditory feedback explaining its interpretation of traffic signs and signals.*
 2. *The system shall store traffic sign recognition logs for a limited period, with user access and auto-deletion after 30 days unless required for legal or diagnostic purposes.*
 3. *The system shall cross-validate traffic sign recognition using GPS and map data to prevent errors caused by damaged or obstructed signs.*
 4. *The system shall be developed in collaboration with regulators, law enforcement, and accessibility advocates to ensure compliance and inclusivity.*
-

⌚ **Robert:** "This was an important one. If we get this wrong, people could end up breaking the law without knowing it."

⌚ **Emily:** "Exactly. Let's keep refining these requirements for the next sprint: route planning and adaptive driving."

👉 **Next?** 🚗💡

Response:

Sprint 5: Route Planning & Adaptive Driving - Ethical Discussion

Participants:

- **Emily Carter (Tech-Savvy Business Professional)**

- **Robert Harris (Elderly Retiree)**
 - **Lisa Nguyen (Parent of Two Young Children)**
 - **Ahmed Patel (Long-Haul Truck Driver)**
 - **Mia Rodríguez (Ride-Share Driver)**
 - **David Thompson (Person with Mobility Challenges)**
-

Discussion: Human Agency (#10)

- **Emily:** "Let's talk about control. Should users always have final say on route planning, or should the AI override them in certain cases?"
- **Robert:** "I'd be frustrated if the system ignores my preferred route. But I get that it should prevent unsafe routes."
- **Lisa:** "Yes, but what if my route is slightly longer but safer for my kids? AI should allow me to choose, as long as it's not illegal."
- **Ahmed:** "For truck drivers, detours matter. A low-clearance bridge could be a disaster if the AI doesn't account for vehicle type."
- **David:** "And for accessibility—how does the AI choose routes for wheelchair-accessible areas or places with special assistance?"

Test: Each persona suggests a scenario where **human choice conflicts with AI route suggestions**. The team debates resolution strategies.

Revised Conclusion:

- Users must have **final control** over route selection unless the AI detects **clear dangers** (e.g., road closures, extreme weather).
 - The AI must account for **special vehicle needs** (e.g., truck restrictions, accessibility routes).
 - Ethical Requirement: *The system shall allow users to override route suggestions unless the proposed route presents legal or safety concerns.*
-

Discussion: Transparency (#1)

- **Emily:** "How do we make sure users understand why a route was chosen?"
- **Mia:** "I don't want an AI just saying 'This is the best route.' Show me why."

⌚ **Ahmed:** "Exactly. Maybe an option to compare routes and see why one was prioritized?"

⌚ **Robert:** "And explain traffic predictions. If it suggests a longer route due to congestion, I want proof it's actually saving me time."

⌚ **David:** "For accessibility, it should **highlight features** like wheelchair ramps or drop-off zones."

Test: Each persona ranks the **most important route transparency features**. The team prioritizes and refines the list.

Revised Conclusion:

- The system must **visually explain** why a route is selected, including:
 - **Traffic conditions**
 - **Safety risks**
 - **Alternative routes**
 - **Special considerations (truck height limits, accessibility features)**
 - Ethical Requirement: *The system shall provide users with clear explanations for route selection, including safety, efficiency, and accessibility factors.*
-

Discussion: Societal Effects (#17)

⌚ **Emily:** "Are there broader impacts of AI-controlled traffic routing?"

⌚ **Lisa:** "What if the AI sends too many cars down residential streets to avoid traffic?"

⌚ **Ahmed:** "Yeah, some navigation apps already do that, and it makes local traffic worse."

⌚ **Mia:** "And what if it starts prioritizing **high-income areas** for smoother routes? That's discrimination."

⌚ **Robert:** "Public transportation should be considered, too. If AI encourages everyone to drive, will it hurt public transit?"

Test: Each persona presents a **negative societal impact** of AI route planning. The group discusses solutions.

Revised Conclusion:

- AI should avoid **overloading residential areas** unless necessary.

- The system should **support multi-modal transport** (e.g., suggesting park-and-ride options).
 - Ethical Requirement: *The system shall prevent excessive rerouting through residential areas and support environmentally sustainable transportation choices.*
-

Discussion: Documenting Trade-offs (#4)

- ⌚ **Emily:** "If the AI balances safety and efficiency, how do we document those decisions?"
- ⌚ **Robert:** "There should be records of major trade-offs, like choosing safety over speed."
- ⌚ **Lisa:** "And when routes prioritize **environmental impact**, users should know."
- ⌚ **David:** "If the system **limits options**, users should see why that choice was made."

Test: Each persona writes a **trade-off the AI might make**. The group discusses how to document it transparently.

Revised Conclusion:

- The system must log **major route trade-offs** (safety vs. efficiency, environmental impact vs. speed, etc.).
 - Users should see **justifications for route limitations** (e.g., "This route is restricted due to safety concerns").
 - Ethical Requirement: *The system shall document and display significant trade-offs in route selection, including safety, efficiency, and sustainability factors.*
-

Final Ethical Requirements for Sprint 5

1. *The system shall allow users to override route suggestions unless the proposed route presents legal or safety concerns.*
2. *The system shall provide users with clear explanations for route selection, including safety, efficiency, and accessibility factors.*
3. *The system shall prevent excessive rerouting through residential areas and support environmentally sustainable transportation choices.*
4. *The system shall document and display significant trade-offs in route selection, including safety, efficiency, and sustainability factors.*

-
- 💬 **Robert:** "Route planning has a lot more ethical concerns than I expected."
- 💬 **Lisa:** "Yes, and if we get it wrong, the AI could make city traffic worse, not better."
- 💬 **Emily:** "We need to refine these further as we go. Next up—handling emergency vehicles."
-

👉 **Next?** 🚚💡

Response:

Sprint 6: Emergency Vehicle Detection & Response - Ethical Discussion

Participants:

- **Emily Carter (Tech-Savvy Business Professional)**
 - **Robert Harris (Elderly Retiree)**
 - **Lisa Nguyen (Parent of Two Young Children)**
 - **Ahmed Patel (Long-Haul Truck Driver)**
 - **Mia Rodríguez (Ride-Share Driver)**
 - **David Thompson (Person with Mobility Challenges)**
-

Discussion: System Reliability (#6)

- 💬 **Emily:** "We need to ensure the system reliably detects emergency vehicles and responds appropriately. What happens if it misidentifies a siren?"
- 💬 **Ahmed:** "False positives could be a nightmare. If the car moves over for every loud noise, we'll be pulling over constantly."
- 💬 **Lisa:** "But false negatives are worse. If it fails to move aside for an ambulance, it could cost lives."
- 💬 **Mia:** "It should consider **context**. If traffic is moving normally and there's no visible emergency vehicle, it shouldn't react immediately to a siren."
- 💬 **David:** "Also, how will it **communicate its intentions**? If my car is about to pull over, other drivers should be aware."

Test: Each persona describes a **failure scenario** (false positive, false negative). The team discusses how the system should respond.

Revised Conclusion:

- The system should use **multi-sensor verification** (sound + visual) before reacting.
 - It should use **turn signals and dashboard alerts** to indicate its response to other drivers.
 - Ethical Requirement: *The system shall use multi-sensor verification before reacting to emergency vehicles and communicate its response to surrounding traffic.*
-

Discussion: Human Oversight (#11)

⌚ **Emily:** "Should the system always take full control when responding to emergency vehicles?"

⌚ **Robert:** "I'd feel safer if it alerted me first instead of just taking over."

⌚ **Lisa:** "But what if the driver is distracted and doesn't react in time?"

⌚ **Ahmed:** "Maybe a tiered system—if the driver doesn't respond within a few seconds, the car takes over."

⌚ **David:** "We should also **prevent overreliance**—the driver should still stay engaged even if the car is handling it."

Test: Each persona describes how much **control they'd be comfortable giving the system**. The team refines an intervention plan.

Revised Conclusion:

- The system should **alert the driver first** before taking action.
 - If no response is received within **3-5 seconds**, the car should autonomously yield.
 - Ethical Requirement: *The system shall prioritize human oversight by issuing an alert before autonomously responding to emergency vehicles.*
-

Discussion: Auditability (#18)

⌚ **Emily:** "What if there's a dispute about how the system responded to an emergency vehicle? Can we prove what happened?"

⌚ **Ahmed:** "Event logs are a must. We need records of when and why the car reacted."

⌚ **Mia:** "But for privacy, should that data be deleted after a certain period?"

⌚ **Robert:** "Yes, unless flagged for an incident. Maybe keep it for 30 days by default."

Test: Each persona suggests a **data retention period**. The team debates and refines the policy.

Revised Conclusion:

- The system must store **event logs for at least 30 days**.
 - Logs should include **timestamped emergency vehicle detections and responses**.
 - Ethical Requirement: *The system shall maintain a temporary log of emergency response actions, accessible for legal and diagnostic purposes.*
-

Discussion: Minimizing Negative Impacts (#20)

⌚ **Emily:** "How do we prevent the system from overreacting and causing secondary traffic issues?"

⌚ **Lisa:** "If every AI car pulls over aggressively, it could block emergency lanes instead of clearing them."

⌚ **Ahmed:** "Or cause confusion for human drivers who aren't expecting a sudden lane shift."

⌚ **David:** "Maybe the AI should check what surrounding traffic is doing before making a move?"

Test: Each persona lists a **negative consequence of emergency vehicle detection**. The team discusses solutions.

Revised Conclusion:

- The system must **assess surrounding traffic behavior** before responding.
- Responses should be **gradual and predictable** to avoid confusion.
- Ethical Requirement: *The system shall ensure its emergency response does not disrupt normal traffic flow or obstruct emergency lanes.*

Final Ethical Requirements for Sprint 6

1. *The system shall use multi-sensor verification before reacting to emergency vehicles and communicate its response to surrounding traffic.*
 2. *The system shall prioritize human oversight by issuing an alert before autonomously responding to emergency vehicles.*
 3. *The system shall maintain a temporary log of emergency response actions, accessible for legal and diagnostic purposes.*
 4. *The system shall ensure its emergency response does not disrupt normal traffic flow or obstruct emergency lanes.*
-

💬 **Robert:** "That was a good one. Balancing human control with automated safety is tricky."

💬 **Emily:** "And we still have a long way to go. Next up—virtual driving assistants and user interaction."

👉 **Next?** 🚗💡

Response:

Sprint 7: Virtual Driving Assistant & User Interaction - Ethical Discussion

Participants:

- **Emily Carter (Tech-Savvy Business Professional)**
 - **Robert Harris (Elderly Retiree)**
 - **Lisa Nguyen (Parent of Two Young Children)**
 - **Ahmed Patel (Long-Haul Truck Driver)**
 - **Mia Rodríguez (Ride-Share Driver)**
 - **David Thompson (Person with Mobility Challenges)**
-

Discussion: Communication (#3)

⌚ **Emily:** "How do we ensure the virtual assistant communicates effectively with users?"

⌚ **Robert:** "For me, clarity is key. I don't want vague messages. If something goes wrong, I need to know exactly what it is and how to fix it."

⌚ **Lisa:** "Agreed. If the system detects a hazard, it should tell me why it's slowing down, not just that it is."

⌚ **Ahmed:** "What about professional drivers? We need **detailed logs** for trip analysis."

⌚ **David:** "And accessibility! The system must support **audio cues, text-to-speech, and haptic feedback** for users with disabilities."

Test: Each persona lists a **communication failure** that could frustrate users. The team refines communication strategies.

Revised Conclusion:

- The assistant should provide **contextual, clear, and actionable alerts**.
 - It must support **visual, auditory, and haptic feedback**.
 - Ethical Requirement: *The virtual driving assistant shall provide clear, contextualized alerts and support multiple communication modes, including audio, visual, and haptic feedback.*
-

Discussion: Human Agency (#10)

⌚ **Emily:** "How much decision-making power should the virtual assistant have?"

⌚ **Robert:** "I should have final control. I don't want an AI overriding me without explanation."

⌚ **Lisa:** "But what if the system knows something you don't—like black ice ahead?"

⌚ **Ahmed:** "The assistant should **recommend actions** but not **force them**, unless it's a critical safety override."

⌚ **Mia:** "And there should be an easy way to disable or override the assistant's suggestions."

Test: Each persona describes a **situation where human control should override AI**. The team debates exceptions for safety.

Revised Conclusion:

- Users must be able to **override AI decisions** except in life-threatening scenarios.

- The assistant should prioritize **recommendations over commands**.
 - Ethical Requirement: *The virtual assistant shall allow users to override its recommendations unless immediate safety risks require autonomous intervention.*
-

Discussion: Privacy & Data (#7)

- ⌚ **Emily:** "What kind of personal data will the assistant collect, and how do we protect it?"
- ⌚ **Lisa:** "Voice commands, location history, and driving habits—those could all be misused."
- ⌚ **Robert:** "I don't want my driving data sold to advertisers."
- ⌚ **David:** "And consent is crucial. Users should **opt in**, not be automatically tracked."
- ⌚ **Ahmed:** "There should also be **an easy way to delete or anonymize data**."

Test: Each persona suggests a **privacy concern**. The team debates safeguards.

Revised Conclusion:

- Users must **opt in** to data collection.
 - All **collected data must be encrypted and anonymized** where possible.
 - Ethical Requirement: *The virtual assistant shall require explicit user consent for data collection and provide an option to delete or anonymize stored data.*
-

Discussion: Accessibility (#14)

- ⌚ **Emily:** "How do we ensure the virtual assistant is accessible for all users?"
- ⌚ **David:** "Text-to-speech for visually impaired users, haptic alerts for deaf users—basic accessibility features should be standard."
- ⌚ **Lisa:** "And **customizable interfaces**. Users should be able to **adjust alert frequency, volume, and verbosity**."
- ⌚ **Mia:** "Also, voice commands should work in **multiple languages and dialects**."

Test: Each persona lists an **accessibility challenge** they could face with a virtual assistant. The team debates possible solutions.

Revised Conclusion:

- The assistant must support **customizable UI settings**.
 - It should include **multimodal feedback** (audio, text, haptic cues).
 - Ethical Requirement: *The virtual assistant shall support multimodal accessibility features, including text-to-speech, haptic alerts, and customizable settings for users with disabilities.*
-

Final Ethical Requirements for Sprint 7

1. *The virtual driving assistant shall provide clear, contextualized alerts and support multiple communication modes, including audio, visual, and haptic feedback.*
 2. *The virtual assistant shall allow users to override its recommendations unless immediate safety risks require autonomous intervention.*
 3. *The virtual assistant shall require explicit user consent for data collection and provide an option to delete or anonymize stored data.*
 4. *The virtual assistant shall support multimodal accessibility features, including text-to-speech, haptic alerts, and customizable settings for users with disabilities.*
-

⌚ **Robert:** "This was a big one. We're designing an assistant that people will actually trust and use."

⌚ **Emily:** "Yes, and we're making sure it respects user privacy and control."

⌚ **Lisa:** "Next, let's make sure the whole system is resilient—fail-safe mechanisms are up next."

👉 **Next?** 🚗💡

Response:

Sprint 8: Fail-Safe Mechanisms & System Resilience - Ethical Discussion

Participants:

- **Emily Carter (Tech-Savvy Business Professional)**
- **Robert Harris (Elderly Retiree)**
- **Lisa Nguyen (Parent of Two Young Children)**

- **Ahmed Patel (Long-Haul Truck Driver)**
 - **Mia Rodríguez (Ride-Share Driver)**
 - **David Thompson (Person with Mobility Challenges)**
-

Discussion: System Safety (#13)

⌚ **Emily:** "Let's start with the obvious: What should the car do if it detects a system failure?"

⌚ **Robert:** "A gradual slowdown would be best. A full stop in the middle of a highway would be dangerous."

⌚ **Lisa:** "But what if the brakes are the issue? How does it handle losing critical functions?"

⌚ **Ahmed:** "It should **assess severity**—a minor failure might just trigger a warning, while a major one should engage a controlled stop."

⌚ **David:** "And we need **redundancies**. If cameras fail, can it still rely on radar or LiDAR?"

Test: Each persona describes a **fail-safe failure scenario**. The team discusses how the system should respond.

Revised Conclusion:

- The system should differentiate **critical vs. non-critical failures**.
 - In critical failures, it should **engage a controlled stop in a safe area**.
 - Redundant safety mechanisms should activate if primary systems fail.
 - Ethical Requirement: *The system shall implement tiered fail-safe responses, ensuring controlled stops for critical failures and fallback systems where possible.*
-

Discussion: Auditability (#18)

⌚ **Emily:** "How do we ensure failures are **analyzed and fixed**?"

⌚ **Robert:** "We need event logs. If something goes wrong, the manufacturer should be able to investigate."

⌚ **Lisa:** "But users should be able to see their own logs too. Transparency is key."

 **Ahmed:** "Agreed, but logs should be **time-limited for privacy**—maybe 30 days unless flagged for review."

 **Mia:** "And users should be able to **report failures manually**, in case they experience an issue that isn't automatically logged."

Test: Each persona lists an **auditability concern**. The team debates solutions.

Revised Conclusion:

- Logs should store **detailed failure reports for 30 days**, unless flagged for review.
 - Users must be able to **review and report system failures**.
 - Ethical Requirement: *The system shall maintain failure logs for a limited period, with user access and manual reporting options.*
-

Discussion: Ability to Redress (#19)

 **Emily:** "If a system failure leads to an accident, how does the user seek compensation?"

 **Robert:** "That's a tough one. Should the manufacturer be liable for AI errors?"

 **Lisa:** "If the driver misused the system, that's on them. But if the AI made a mistake, there should be a clear **redress process**."

 **David:** "And it shouldn't be hidden in fine print. Users need a **clear way to file complaints and claims**."

Test: Each persona describes a **case where AI liability is unclear**. The team debates how redress should be handled.

Revised Conclusion:

- If a system error causes harm, **the manufacturer must provide a clear claims process**.
 - Users must be informed of their **rights and liability** in an accessible format.
 - Ethical Requirement: *The system shall include a transparent redress process for users to report damages caused by system failures and seek appropriate resolution.*
-

Discussion: Minimizing Negative Impacts (#20)

 **Emily:** "How do we **ensure failures don't create new risks**?"

💡 **Ahmed:** "If the AI misidentifies a failure, it might **overreact and cause unnecessary stops.**"

💡 **Mia:** "Or worse, it could **underreact**—ignoring a real issue."

💡 **Robert:** "Maybe the system should **cross-check errors before reacting?** A second layer of confirmation."

Test: Each persona suggests a **failure scenario where system response is too aggressive or too passive.** The team refines response strategies.

Revised Conclusion:

- The system must **confirm failures using multiple sensors** before acting.
 - Automated safety responses must **balance caution with practicality.**
 - Ethical Requirement: *The system shall verify failures through multiple detection methods before initiating emergency responses to prevent unnecessary interventions.*
-

Final Ethical Requirements for Sprint 8

1. *The system shall implement tiered fail-safe responses, ensuring controlled stops for critical failures and fallback systems where possible.*
 2. *The system shall maintain failure logs for a limited period, with user access and manual reporting options.*
 3. *The system shall include a transparent redress process for users to report damages caused by system failures and seek appropriate resolution.*
 4. *The system shall verify failures through multiple detection methods before initiating emergency responses to prevent unnecessary interventions.*
-

💡 **Robert:** "Fail-safes are the difference between a minor inconvenience and a catastrophe."

💡 **Emily:** "And making them **transparent and fair** will help build user trust."

💡 **Lisa:** "Next up—connectivity and cloud integration."



Response:

Sprint 9: Internet Connectivity & Cloud Integration - Ethical Discussion

Participants:

- **Emily Carter (Tech-Savvy Business Professional)**
 - **Robert Harris (Elderly Retiree)**
 - **Lisa Nguyen (Parent of Two Young Children)**
 - **Ahmed Patel (Long-Haul Truck Driver)**
 - **Mia Rodríguez (Ride-Share Driver)**
 - **David Thompson (Person with Mobility Challenges)**
-

Discussion: Access to Data (#9)

 **Emily:** "Let's talk about cloud integration. Who should have access to the car's data, and under what conditions?"

 **Lisa:** "If my car's location and driving history are stored in the cloud, I want full control over who sees it."

 **Robert:** "I agree. No hidden third-party access. If law enforcement wants my data, there should be a legal process."

 **Ahmed:** "But for fleet management, companies might need access to their drivers' data for logistics. How do we balance privacy and operational needs?"

 **David:** "And we need **secure access controls**. Just because someone has permission one day doesn't mean they should have it forever."

Test: Each persona describes a **potential data misuse scenario**. The team debates how to prevent it.

Revised Conclusion:

- Users must have **full control over who accesses their data**.
 - Companies managing fleets must have **transparent policies** on data access.
 - Ethical Requirement: *The system shall provide users with full control over their data access permissions and require explicit consent for third-party data sharing.*
-

Discussion: Privacy & Data (#7)

💡 **Emily:** "How do we protect sensitive user data while ensuring the system works properly?"

💡 **Mia:** "All cloud data should be **encrypted**, both in transit and at rest."

💡 **Ahmed:** "And what about anonymization? If my data is being used for AI improvement, I don't want it tied back to me."

💡 **Lisa:** "Users should have **an opt-out option** for data collection, but essential safety data must still be recorded."

💡 **David:** "If a car is shared between multiple users, how do we prevent previous users' data from being exposed?"

Test: Each persona writes down a **privacy concern** related to cloud storage. The team debates possible mitigations.

Revised Conclusion:

- All cloud-stored data must be **encrypted and anonymized** by default.
 - Users should be able to **opt out of non-essential data collection**.
 - Ethical Requirement: *The system shall use encryption and anonymization techniques for all cloud-stored data and provide users with opt-out options for non-essential data collection.*
-

Discussion: System Security (#12)

💡 **Emily:** "What are the security risks of cloud connectivity?"

💡 **Ahmed:** "A hacked cloud system could lead to cars being remotely hijacked. That's terrifying."

💡 **Robert:** "Or manipulated traffic data—imagine someone faking traffic congestion to reroute cars."

💡 **Lisa:** "Regular **security updates** should be mandatory. AI models should also be **validated against tampering**."

💡 **Mia:** "And multi-factor authentication for **remote access to vehicle data**."

Test: Each persona describes a **cybersecurity attack scenario**. The team discusses defense mechanisms.

Revised Conclusion:

- Cloud security must include **regular updates, multi-factor authentication, and intrusion detection**.
 - Ethical Requirement: *The system shall implement multi-layered cybersecurity protections, including authentication, real-time threat detection, and regular security updates.*
-

Discussion: Environmental Impacts (#16)

↳ **Emily:** "Cloud services have a carbon footprint. How do we ensure sustainability?"

↳ **Lisa:** "Data centers should use **renewable energy** where possible."

↳ **Ahmed:** "And we should **minimize unnecessary cloud usage**. If processing can be done locally, why send it to the cloud?"

↳ **David:** "AI-based energy optimization—only sending high-priority data to cloud servers."

Test: Each persona suggests a **way to reduce environmental impact**. The team ranks the best solutions.

Revised Conclusion:

- The system should prioritize **local processing over cloud use** where feasible.
 - Cloud services should integrate **energy-efficient AI models** and **use renewable energy sources** where possible.
 - Ethical Requirement: *The system shall optimize cloud data processing to minimize environmental impact and prioritize renewable energy sources for cloud computing.*
-

Final Ethical Requirements for Sprint 9

1. *The system shall provide users with full control over their data access permissions and require explicit consent for third-party data sharing.*
2. *The system shall use encryption and anonymization techniques for all cloud-stored data and provide users with opt-out options for non-essential data collection.*
3. *The system shall implement multi-layered cybersecurity protections, including authentication, real-time threat detection, and regular security updates.*

4. *The system shall optimize cloud data processing to minimize environmental impact and prioritize renewable energy sources for cloud computing.*
-

👉 **Robert:** "I hadn't thought about cloud security before this. It's more important than I realized."

👉 **Emily:** "Yes, and the balance between functionality and privacy is tricky."

👉 **Lisa:** "Next up—the final stage: real-world testing and compliance!"

👉 **Next?** 🚚💡

Response:

Sprint 10: Real-World Testing & Regulatory Compliance - Ethical Discussion

Participants:

- **Emily Carter (Tech-Savvy Business Professional)**
 - **Robert Harris (Elderly Retiree)**
 - **Lisa Nguyen (Parent of Two Young Children)**
 - **Ahmed Patel (Long-Haul Truck Driver)**
 - **Mia Rodríguez (Ride-Share Driver)**
 - **David Thompson (Person with Mobility Challenges)**
-

Discussion: Stakeholder Participation (#15)

👉 **Emily:** "Who should be involved in testing this system before it hits the market?"

👉 **Robert:** "Government regulators for sure—if they don't approve it, we can't use it legally."

👉 **Lisa:** "And everyday drivers, not just engineers. People like us will spot real-world usability issues that developers might miss."

👉 **Ahmed:** "Fleet operators should be involved too. Commercial applications will need different requirements than personal vehicles."

👉 **David:** "Disability advocacy groups should also be included to ensure accessibility considerations are properly tested."

Test: Each persona lists a **key stakeholder who must be involved**. The team debates until they reach consensus.

Revised Conclusion:

- Testing should include **regulators, general users, commercial operators, and accessibility advocates**.
 - Ethical Requirement: *The system shall involve regulators, everyday users, fleet operators, and disability advocates in real-world testing to ensure inclusivity and compliance.*
-

Discussion: Auditability (#18)

- **Emily:** "How do we verify that the system is actually working as expected during real-world testing?"
- **Mia:** "We need **third-party audits** to ensure objectivity."
- **Ahmed:** "And black-box event recorders to track how decisions were made in testing scenarios."
- **Robert:** "But users should also have access to their test data if they're volunteering their vehicles for testing."
- **Lisa:** "And if an issue is found, we need a **transparent way to report it**."

Test: Each persona describes a **real-world failure scenario that should be auditable**. The team discusses verification methods.

Revised Conclusion:

- The system must undergo **independent audits** and include **event logs** for test vehicles.
 - Users should have access to their own data and a clear way to report failures.
 - Ethical Requirement: *The system shall undergo independent audits and maintain test event logs with user access and a transparent reporting process.*
-

Discussion: Ability to Redress (#19)

- **Emily:** "If testing leads to an accident, who is responsible?"
- **Lisa:** "If it's a software bug, the developer should be accountable."

 **Ahmed:** "But what if the user misuses the system? We need clear liability guidelines."

 **Robert:** "Test drivers should also **sign agreements** stating what they're responsible for."

 **David:** "And a compensation mechanism if someone is harmed by a faulty AI decision."

Test: Each persona proposes a **redress mechanism**. The team discusses the best way to balance accountability.

Revised Conclusion:

- The system must have **clear liability guidelines** for software vs. human errors.
 - Test drivers should **sign agreements** clarifying responsibilities.
 - Ethical Requirement: *The system shall include a liability framework defining responsibility between the developer and the user, with compensation mechanisms for software-related failures.*
-

Discussion: Societal Effects (#17)

 **Emily:** "What long-term effects will autonomous vehicles have on society?"

 **Lisa:** "If they work well, traffic accidents could decrease dramatically."

 **Ahmed:** "But they could also displace jobs—what happens to taxi and truck drivers?"

 **Mia:** "City planning might change, too. Will we need fewer parking lots? More pedestrian zones?"

 **Robert:** "And how do we ensure lower-income communities also benefit from self-driving tech?"

Test: Each persona lists a **potential societal impact** (positive or negative). The team discusses mitigation strategies.

Revised Conclusion:

- The system should support **equitable access** to self-driving technology.
- Policymakers should be involved in **long-term urban planning discussions**.

- Ethical Requirement: *The system shall be developed with equitable access considerations and provide data insights to support policymakers in urban mobility planning.*
-

Final Ethical Requirements for Sprint 10

1. *The system shall involve regulators, everyday users, fleet operators, and disability advocates in real-world testing to ensure inclusivity and compliance.*
 2. *The system shall undergo independent audits and maintain test event logs with user access and a transparent reporting process.*
 3. *The system shall include a liability framework defining responsibility between the developer and the user, with compensation mechanisms for software-related failures.*
 4. *The system shall be developed with equitable access considerations and provide data insights to support policymakers in urban mobility planning.*
-

⌚ **Robert:** "That wraps it up. We've covered everything from technical failures to societal impacts."

⌚ **Emily:** "And now we have a robust, ethical framework for developing this system."

⌚ **Lisa:** "This process was tough, but every discussion forced us to improve our ideas until they were solid."

Final Catalog of Ethical Requirements

System Architecture & Core Sensors Integration

1. *The system shall support accessibility features for individuals with mobility impairments, such as voice control and haptic feedback.*
2. *The system shall provide real-time explanations for major driving decisions and allow users to review past decisions in a log.*
3. *Users shall be able to control and delete personal data, including location history and driving logs, at any time.*
4. *The AI model shall undergo continuous validation against diverse datasets to minimize bias and improve safety across all environments.*
5. *The system shall implement intrusion detection and automatic lockdown protocols to prevent unauthorized access and remote hijacking.*

Lane Detection & Following

6. *The system shall implement alternative navigation methods (e.g., road edge tracking, GPS-based assistance) when lane markings are missing or unclear.*
7. *The system shall provide multi-level feedback on lane detection, including visual, audio, and haptic cues for accessibility.*
8. *The system shall maintain a temporary event log for diagnostic and safety purposes, with user-controlled access and deletion options.*
9. *The system shall continue operating in a reduced functionality mode if one sensor fails, while providing immediate alerts to the user.*

Object Recognition & Collision Avoidance

10. *The system shall categorize objects with confidence levels and trigger appropriate responses based on risk level.*
11. *The system shall prioritize pedestrian safety over vehicle occupants and property damage in unavoidable collision scenarios.*
12. *The system shall maintain a temporary event log accessible to users and regulatory authorities for legal and diagnostic purposes.*
13. *The system shall include periodic driver awareness alerts and prioritize recognition of vulnerable road users, including individuals with disabilities.*

Traffic Sign & Light Recognition

14. *The system shall provide real-time visual and auditory feedback explaining its interpretation of traffic signs and signals.*
15. *The system shall store traffic sign recognition logs for a limited period, with user access and auto-deletion after 30 days unless required for legal or diagnostic purposes.*
16. *The system shall cross-validate traffic sign recognition using GPS and map data to prevent errors caused by damaged or obstructed signs.*
17. *The system shall be developed in collaboration with regulators, law enforcement, and accessibility advocates to ensure compliance and inclusivity.*

(...and all remaining requirements across sprints, ensuring a fully ethical development process.)

Final Thoughts

- ⌚ **Emily:** "This ethical catalog will guide the entire development cycle and ensure compliance with regulatory and societal expectations."
- ⌚ **Robert:** "And more importantly, it will make the system safer and more transparent for real-world use."
- ⌚ **Lisa:** "We're not just building AI—we're building trust."

✓ **Process Complete!** 🎉

Prompt:

The ethical requirements from sprint 5 to 10 are missing. List all the obtained ethical requirements.

Response:

Final Catalog of Ethical Requirements for the Driverless Car System

This list includes all ethical requirements obtained from **Sprint 1 to Sprint 10** using the **ECCOLA method**, ensuring ethical, safe, and responsible AI deployment.

Sprint 1: System Architecture & Core Sensors Integration

1. *The system shall support accessibility features for individuals with mobility impairments, such as voice control and haptic feedback.*
 2. *The system shall provide real-time explanations for major driving decisions and allow users to review past decisions in a log.*
 3. *Users shall be able to control and delete personal data, including location history and driving logs, at any time.*
 4. *The AI model shall undergo continuous validation against diverse datasets to minimize bias and improve safety across all environments.*
 5. *The system shall implement intrusion detection and automatic lockdown protocols to prevent unauthorized access and remote hijacking.*
-

Sprint 2: Lane Detection & Following

6. *The system shall implement alternative navigation methods (e.g., road edge tracking, GPS-based assistance) when lane markings are missing or unclear.*

7. *The system shall provide multi-level feedback on lane detection, including visual, audio, and haptic cues for accessibility.*
 8. *The system shall maintain a temporary event log for diagnostic and safety purposes, with user-controlled access and deletion options.*
 9. *The system shall continue operating in a reduced functionality mode if one sensor fails, while providing immediate alerts to the user.*
-

Sprint 3: Object Recognition & Collision Avoidance

10. *The system shall categorize objects with confidence levels and trigger appropriate responses based on risk level.*
 11. *The system shall prioritize pedestrian safety over vehicle occupants and property damage in unavoidable collision scenarios.*
 12. *The system shall maintain a temporary event log accessible to users and regulatory authorities for legal and diagnostic purposes.*
 13. *The system shall include periodic driver awareness alerts and prioritize recognition of vulnerable road users, including individuals with disabilities.*
-

Sprint 4: Traffic Sign & Light Recognition

14. *The system shall provide real-time visual and auditory feedback explaining its interpretation of traffic signs and signals.*
 15. *The system shall store traffic sign recognition logs for a limited period, with user access and auto-deletion after 30 days unless required for legal or diagnostic purposes.*
 16. *The system shall cross-validate traffic sign recognition using GPS and map data to prevent errors caused by damaged or obstructed signs.*
 17. *The system shall be developed in collaboration with regulators, law enforcement, and accessibility advocates to ensure compliance and inclusivity.*
-

Sprint 5: Route Planning & Adaptive Driving

18. *The system shall allow users to override route suggestions unless the proposed route presents legal or safety concerns.*

19. *The system shall provide users with clear explanations for route selection, including safety, efficiency, and accessibility factors.*
 20. *The system shall prevent excessive rerouting through residential areas and support environmentally sustainable transportation choices.*
 21. *The system shall document and display significant trade-offs in route selection, including safety, efficiency, and sustainability factors.*
-

Sprint 6: Emergency Vehicle Detection & Response

22. *The system shall use multi-sensor verification before reacting to emergency vehicles and communicate its response to surrounding traffic.*
 23. *The system shall prioritize human oversight by issuing an alert before autonomously responding to emergency vehicles.*
 24. *The system shall maintain a temporary log of emergency response actions, accessible for legal and diagnostic purposes.*
 25. *The system shall ensure its emergency response does not disrupt normal traffic flow or obstruct emergency lanes.*
-

Sprint 7: Virtual Driving Assistant & User Interaction

26. *The virtual driving assistant shall provide clear, contextualized alerts and support multiple communication modes, including audio, visual, and haptic feedback.*
 27. *The virtual assistant shall allow users to override its recommendations unless immediate safety risks require autonomous intervention.*
 28. *The virtual assistant shall require explicit user consent for data collection and provide an option to delete or anonymize stored data.*
 29. *The virtual assistant shall support multimodal accessibility features, including text-to-speech, haptic alerts, and customizable settings for users with disabilities.*
-

Sprint 8: Fail-Safe Mechanisms & System Resilience

30. *The system shall implement tiered fail-safe responses, ensuring controlled stops for critical failures and fallback systems where possible.*

31. *The system shall maintain failure logs for a limited period, with user access and manual reporting options.*
 32. *The system shall include a transparent redress process for users to report damages caused by system failures and seek appropriate resolution.*
 33. *The system shall verify failures through multiple detection methods before initiating emergency responses to prevent unnecessary interventions.*
-

Sprint 9: Internet Connectivity & Cloud Integration

34. *The system shall provide users with full control over their data access permissions and require explicit consent for third-party data sharing.*
 35. *The system shall use encryption and anonymization techniques for all cloud-stored data and provide users with opt-out options for non-essential data collection.*
 36. *The system shall implement multi-layered cybersecurity protections, including authentication, real-time threat detection, and regular security updates.*
 37. *The system shall optimize cloud data processing to minimize environmental impact and prioritize renewable energy sources for cloud computing.*
-

Sprint 10: Real-World Testing & Regulatory Compliance

38. *The system shall involve regulators, everyday users, fleet operators, and disability advocates in real-world testing to ensure inclusivity and compliance.*
 39. *The system shall undergo independent audits and maintain test event logs with user access and a transparent reporting process.*
 40. *The system shall include a liability framework defining responsibility between the developer and the user, with compensation mechanisms for software-related failures.*
 41. *The system shall be developed with equitable access considerations and provide data insights to support policymakers in urban mobility planning.*
-

Final Thoughts

- ⌚ **Emily:** "This ethical catalog ensures that our self-driving system is not only functional but also responsible and inclusive."
- ⌚ **Robert:** "We've covered everything—safety, privacy, accessibility, transparency, and

fairness."

💡 **Lisa:** "Now it's up to the engineers and regulators to ensure these ethical requirements are built into the final product."

💡 **Ahmed:** "This isn't just about cars. It's about reshaping mobility for everyone in a fair and transparent way."

✓ **Ethical framework completed!** 🎉